



ALVAREZ

College of Business

The University of Texas at San Antonio

Introduction to Programming in R

Module 5:
Data Wrangling
Part 2

Learning Objectives

- Convert data from long to short format or vice versa.
- Separate data from one column to multiple.
- Unite multiple columns of data into one.
- Relational data, using keys to join datasets.

Pivot: Long

- Often data may look tidy at first glance, but it is not. For e.g., some observations may be stored as variables.
- For e.g. take this built-in data on religion and income.
- Do `data("relig_income")` and `print(relig_income)`.

```
> relig_income
# A tibble: 18 x 11
  religion `<$10k` `<$10-20k` `<$20-30k` `<$30-40k` `<$40-50k` `<$50-75k` `<$75-100k` `<$100-150k` `<150k` `Don't know/ref...
  <chr>      <dbl>      <dbl>      <dbl>      <dbl>      <dbl>      <dbl>      <dbl>      <dbl>      <dbl>      <dbl>
1 Agnostic    27         34         60         81         76        137        122        109         84         96
2 Atheist     12         27         37         52         35         70         73         59         74         76
3 Buddhist    27         21         30         34         33         58         62         39         53         54
4 Catholic   418        617        732        670        638       1116        949        792        633       1489
5 Don't kn...  15         14         15         11         10         35         21         17         18        116
6 Evangelic... 575        869       1064        982        881       1486        949        723        414       1529
7 Hindu        1          9          7          9         11         34         47         48         54         37
8 Historic... 228        244        236        238        197        223        131         81         78        339
9 Jehovah'...  20         27         24         24         21         30         15         11          6         37
10 Jewish     19         19         25         25         30         95         69         87        151        162
```

- The income levels (variables) are observation data.
- This format of data is referred to as wider format. We can convert from wider to longer format.

Pivot: Long (cont.)

- Long format: Data Variables Name of variable Counts
- ```
> pivot_longer(relig_income, -religion, names_to = "income", values_to = "count")
A tibble: 180 x 3
```

|    | religion | income             | count |
|----|----------|--------------------|-------|
|    | <chr>    | <chr>              | <dbl> |
| 1  | Agnostic | <\$10k             | 27    |
| 2  | Agnostic | \$10-20k           | 34    |
| 3  | Agnostic | \$20-30k           | 60    |
| 4  | Agnostic | \$30-40k           | 81    |
| 5  | Agnostic | \$40-50k           | 76    |
| 6  | Agnostic | \$50-75k           | 137   |
| 7  | Agnostic | \$75-100k          | 122   |
| 8  | Agnostic | \$100-150k         | 109   |
| 9  | Agnostic | >150k              | 84    |
| 10 | Agnostic | Don't know/refused | 96    |

# ... with 170 more rows

The pivot arguments works like *select()*.

Variables to observations

# Pivot: Wide to Long

- Do `data("billboard")` and `print(billboard)`

```
> print(as_tibble(billboard))
A tibble: 317 x 79
 artist track date.entered wk1 wk2 wk3 wk4 wk5 wk6 wk7 wk8 wk9 wk10 wk11 wk12 wk13 wk14 wk15
 <chr> <chr> <date> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
1 2 Pac Baby... 2000-02-26 87 82 72 77 87 94 99 NA NA NA NA NA NA NA NA NA
2 2Ge+h... The ... 2000-09-02 91 87 92 NA NA NA NA NA NA NA NA NA NA NA NA NA
3 3 Doo... Kryp... 2000-04-08 81 70 68 67 66 57 54 53 51 51 51 51 47 44 38
4 3 Doo... Loser 2000-10-21 76 76 72 69 67 65 55 59 62 61 61 59 61 66 72
5 504 B... Wobb... 2000-04-15 57 34 25 17 17 31 36 49 53 57 64 70 75 76 78
6 98^0 Give... 2000-08-19 51 39 34 26 26 19 2 2 3 6 7 22 29 36 47
7 A*Tee... Danc... 2000-07-08 97 97 96 95 100 NA NA NA NA NA NA NA NA NA
8 Aaliy... I Do... 2000-01-29 84 62 51 41 38 35 35 38 38 36 37 37 38 49 61
9 Aaliy... Try ... 2000-03-18 59 53 38 28 21 18 16 14 12 10 9 8 6 1 2
10 Adams... Open... 2000-08-26 76 76 74 69 68 67 61 58 57 59 66 68 61 67 59
... with 307 more rows, and 61 more variables: wk16 <dbl>, wk17 <dbl>, wk18 <dbl>, wk19 <dbl>, wk20 <dbl>,
```

- You want to rank the songs for each week consequently, in a longer format.
- I want include all columns that start with 'wk', remove the prefix 'wk' from each variable, and ignore the NAs.

```
> pivot_longer(billboard, cols = starts_with("wk"), names_to = "week", names_prefix = "wk", values_to = "rank", values_drop_na = TRUE)
```

# Pivot: Long

Ignoring the NAs makes sure we stop at the week with the final ranking info.

```
> pivot_longer(billboard, cols = starts_with("wk"), names_to = "week", names_prefix = "wk", values_to = "rank", values_drop_na = TRUE)
A tibble: 5,307 x 5
 artist track date.entered week rank
 <chr> <chr> <date> <chr> <dbl>
1 2 Pac Baby Don't Cry (Keep... 2000-02-26 1 87
2 2 Pac Baby Don't Cry (Keep... 2000-02-26 2 82
3 2 Pac Baby Don't Cry (Keep... 2000-02-26 3 72
4 2 Pac Baby Don't Cry (Keep... 2000-02-26 4 77
5 2 Pac Baby Don't Cry (Keep... 2000-02-26 5 87
6 2 Pac Baby Don't Cry (Keep... 2000-02-26 6 94
7 2 Pac Baby Don't Cry (Keep... 2000-02-26 7 99
8 2Ge+her The Hardest Part Of ... 2000-09-02 1 91
9 2Ge+her The Hardest Part Of ... 2000-09-02 2 87
10 2Ge+her The Hardest Part Of ... 2000-09-02 3 92
... with 5,297 more rows
```

- Often the data comes in a longer format and we want to convert to wider format.
- You can do this with `pivot_wider()`.
- Load the data `us_rent_income` for a long format example.

`data("us_rent_income")` and `print(us_rent_income)`

What if I want both columns: *est of income*, *est of rent*, *moe of income*, *moe of rent*

```
> print(us_rent_income)
A tibble: 104 x 5
 GEOID NAME variable estimate moe
 <chr> <chr> <chr> <dbl> <dbl>
1 01 Alabama income 24476 136
2 01 Alabama rent 747 3
3 02 Alaska income 32940 508
4 02 Alaska rent 1200 13
5 04 Arizona income 27517 148
6 04 Arizona rent 972 4
7 05 Arkansas income 23789 165
8 05 Arkansas rent 709 5
9 06 California income 29454 109
10 06 California rent 1358 3
... with 94 more rows
```

# Pivot: Long to Wide

- If there are multiple columns in the *values\_from*, the column name will be appended to the front of the names\_from value.
  - estimate\_income, estimate\_rent, moe\_income, moe\_rent.*

```
> us_rent_income
A tibble: 104 x 5
 GEOID NAME variable estimate moe
 <chr> <chr> <chr> <dbl> <dbl>
1 01 Alabama income 24476 136
2 01 Alabama rent 747 3
3 02 Alaska income 32940 508
4 02 Alaska rent 1200 13
5 04 Arizona income 27517 148
6 04 Arizona rent 972 4
7 05 Arkansas income 23789 165
8 05 Arkansas rent 709 5
9 06 California income 29454 109
10 06 California rent 1358 3
... with 94 more rows
> pivot_wider(us_rent_income, names_from = variable, values_from = c(estimate, moe))
A tibble: 52 x 6
 GEOID NAME estimate_income estimate_rent moe_income moe_rent
 <chr> <chr> <dbl> <dbl> <dbl> <dbl>
1 01 Alabama 24476 747 136 3
2 02 Alaska 32940 1200 508 13
3 04 Arizona 27517 972 148 4
4 05 Arkansas 23789 709 165 5
5 06 California 29454 1358 109 3
6 08 Colorado 32401 1125 109 5
7 09 Connecticut 35326 1123 195 5
8 10 Delaware 31560 1076 247 10
9 11 District of Columbia 43198 1424 681 17
10 12 Florida 25952 1077 70 3
... with 42 more rows
```

# Example

1. For *mtcars*, make the mpg value wider by expanding variables for manual and automatic transmissions, as well as, number of cylinders.
2. Reorganize the tibble with the new variables at the front.

Load the data warpbreaks – about breaks in loom (knitting).

3. Make the data wider by expanding the type of wool used.
4. Are there multiple observations for each wool type and tension?
  - Used *values\_fn* to impose mean on the multiple observations.



# *separate()*

- Often, when we read in .txt files or .csv files, a single column may contain data from 2 variables, separated by a delimiter.
- The argument for *separate()* includes the name of the column to separate and the names of the new columns.
- By default, the function will *separate()* at non-alphanumeric character.
  - You can provide a user-specific delimiter using *sep*. For e.g. *sep = ","* will separate at a comma.

- Simple example df =

| x   |
|-----|
| NA  |
| a.b |
| a.d |
| b.c |



```
> separate(df, x, c("A", "B"))
```

|   | A    | B    |
|---|------|------|
| 1 | <NA> | <NA> |
| 2 | a    | b    |
| 3 | a    | d    |
| 4 | b    | c    |

## separate() (cont.)

- By default, `separate()` will format the new column as a character variable.
- When we want alternatives, we can use the argument `convert`.
- Simple example: `df =`

| x   |
|-----|
| NA  |
| a/1 |
| a/2 |
| b/3 |

```
> y = separate(df, x, c("A", "B"), sep = "/", convert = TRUE)
> print(y)
```

|   | A    | B  |
|---|------|----|
| 1 | <NA> | NA |
| 2 | a    | 1  |
| 3 | a    | 2  |
| 4 | b    | 3  |

```
> class(y$B)
[1] "integer"
```

When the separated column is numeric, `sep = positive value` counts from left, negative from right.

- Simple example: `df =`

| x    |
|------|
| 2019 |
| 2020 |
| 2019 |
| 2020 |
| 2020 |
| 2019 |
| 2019 |

```
y = separate(df, x, c("century", "year"), sep = 2, convert = TRUE)
> print(y)
```

|   | century | year |
|---|---------|------|
| 1 | 20      | 19   |
| 2 | 20      | 20   |
| 3 | 20      | 19   |
| 4 | 20      | 20   |
| 5 | 20      | 20   |
| 6 | 20      | 19   |
| 7 | 20      | 19   |

```
> class(y$century)
[1] "integer"
> class(y$year)
[1] "integer"
```

# *unite()*

- Opposite of `separate()`, combines multiple columns.

- Previous simple example df =

|   | century | year |
|---|---------|------|
| 1 | 20      | 19   |
| 2 | 20      | 20   |
| 3 | 20      | 19   |
| 4 | 20      | 20   |
| 5 | 20      | 20   |
| 6 | 20      | 19   |
| 7 | 20      | 19   |



```
> unite(y, fullyear, century, year)
```

|   | fullyear |
|---|----------|
| 1 | 20_19    |
| 2 | 20_20    |
| 3 | 20_19    |
| 4 | 20_20    |
| 5 | 20_20    |
| 6 | 20_19    |
| 7 | 20_19    |

- The arguments of `unite` include the data, the new name of the combined column, followed by the columns to combine.
- By default, the separator is an underscore `'_'`. You can control this by `sep`.

```
> unite(y, fullyear, century, year, sep = "")
```

|   | fullyear |
|---|----------|
| 1 | 2019     |
| 2 | 2020     |
| 3 | 2019     |
| 4 | 2020     |
| 5 | 2020     |
| 6 | 2019     |
| 7 | 2019     |

# Examples

Load the *esoph* data - alcohol, tobacco, and esophagus cancer data.

1. Separate the *agegp* into two variables: *MinAge* and *MaxAge*.
2. Then combine the age ranges back to *agegp* so it looks like original data.
  - How were the 75+ group handled? What can we do about the NAs?
  - You'll be happy to know that as of Dec 2019, *tidyverse* has *na.rm* option in *unite()*.
3. Make sure you get rid of the NAs in the *agegp* column.
4. Finally, make the data wider by expanding the age groups for both *ncases* and *ncontrols*.

# Relational Data

- Often, multiple data frames contain related data and we need to combine and separate data from within these tables. Useful with related databases (SQL).
- Install package - `install.packages("nycflights13")` and load `library(nycflights13)`.
- Contains 5 related tables: flights, airlines, planes, weather, and airports.

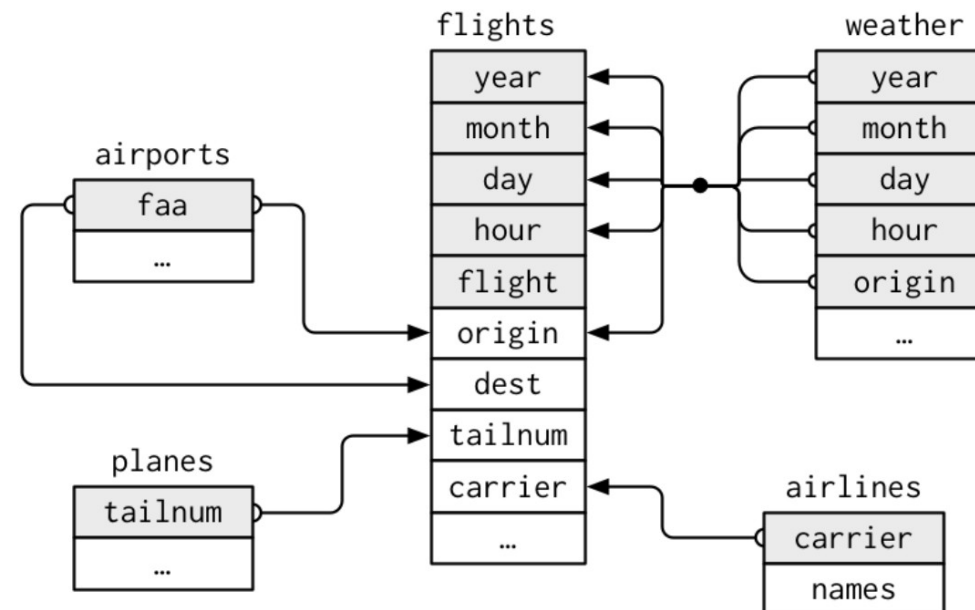


Figure from <https://r4ds.had.co.nz/relational-data.html>

# Relational Data (cont.)

- The variable(s) that connect the tables are called keys. Unique to an observation.
- Two types of keys: primary and foreign.
  - Primary key uniquely identifies an observation in its own table. For e.g. *planes\$tailnum* uniquely identifies each observation in the planes table.
  - Foreign key uniquely identifies an observation in another table. For e.g. *flights\$tailnum* uniquely identifies observation in planes table.
  - We can have variables that is both primary and foreign. For e.g. *weather\$origin* is primary in the weather table but also uniquely identifies observations in airports table.
- Always good to double check if the primary keys are unique. Use *count()* & *filter()*.

```
> dist = count(planes, planes$tailnum)
> print(dist)
A tibble: 3,322 x 2
 `planes$tailnum` n
 <chr> <int>
1 N10156 1
2 N102UW 1
3 N103US 1
4 N104UW 1
5 N10575 1
6 N105UW 1
7 N107US 1
8 N108UW 1
9 N109UW 1
10 N110UW 1
... with 3,312 more rows
```

```
> filter(dist, n > 1)
A tibble: 0 x 2
... with 2 variables: `planes$tailnum` <chr>, n <int>
```

# Relational Data (cont..)

- Sometime, tables don't have primary keys. For e.g. the flights table does not have one. What about *year*, *month*, *day*, and *flight*?

- For such cases, you can use

`mutate(flights, id = row_number())`

```
> dist = count(flights, flights$year, flights$month, flights$day, flights$flight)
> filter(dist, n > 1)
A tibble: 29,768 x 5
 `flights$year` `flights$month` `flights$day` `flights$flight` n
 <int> <int> <int> <int> <int>
1 2013 1 1 1 2
2 2013 1 1 3 2
3 2013 1 1 4 2
4 2013 1 1 11 3
5 2013 1 1 15 2
6 2013 1 1 21 2
7 2013 1 1 27 4
8 2013 1 1 31 2
9 2013 1 1 32 2
10 2013 1 1 35 2
... with 29,758 more rows
```

| hour | temp  | dewp  | humid | wind_dir | wind_speed | wind_gust | precip | pressure | visib | time_hour           | id |
|------|-------|-------|-------|----------|------------|-----------|--------|----------|-------|---------------------|----|
| 1    | 39.02 | 26.06 | 59.37 | 270      | 10.35702   | NA        | 0      | 1012.0   | 10    | 2013-01-01 01:00:00 | 1  |
| 2    | 39.02 | 26.96 | 61.63 | 250      | 8.05546    | NA        | 0      | 1012.3   | 10    | 2013-01-01 02:00:00 | 2  |
| 3    | 39.02 | 28.04 | 64.43 | 240      | 11.50780   | NA        | 0      | 1012.5   | 10    | 2013-01-01 03:00:00 | 3  |
| 4    | 39.92 | 28.04 | 62.21 | 250      | 12.65858   | NA        | 0      | 1012.2   | 10    | 2013-01-01 04:00:00 | 4  |
| 5    | 39.02 | 28.04 | 64.43 | 260      | 12.65858   | NA        | 0      | 1011.9   | 10    | 2013-01-01 05:00:00 | 5  |
| 6    | 37.94 | 28.04 | 67.21 | 240      | 11.50780   | NA        | 0      | 1012.4   | 10    | 2013-01-01 06:00:00 | 6  |
| 7    | 39.02 | 28.04 | 64.43 | 240      | 14.96014   | NA        | 0      | 1012.2   | 10    | 2013-01-01 07:00:00 | 7  |
| 8    | 39.92 | 28.04 | 62.21 | 250      | 10.35702   | NA        | 0      | 1012.2   | 10    | 2013-01-01 08:00:00 | 8  |

# Combining Tables: Mutating Joins

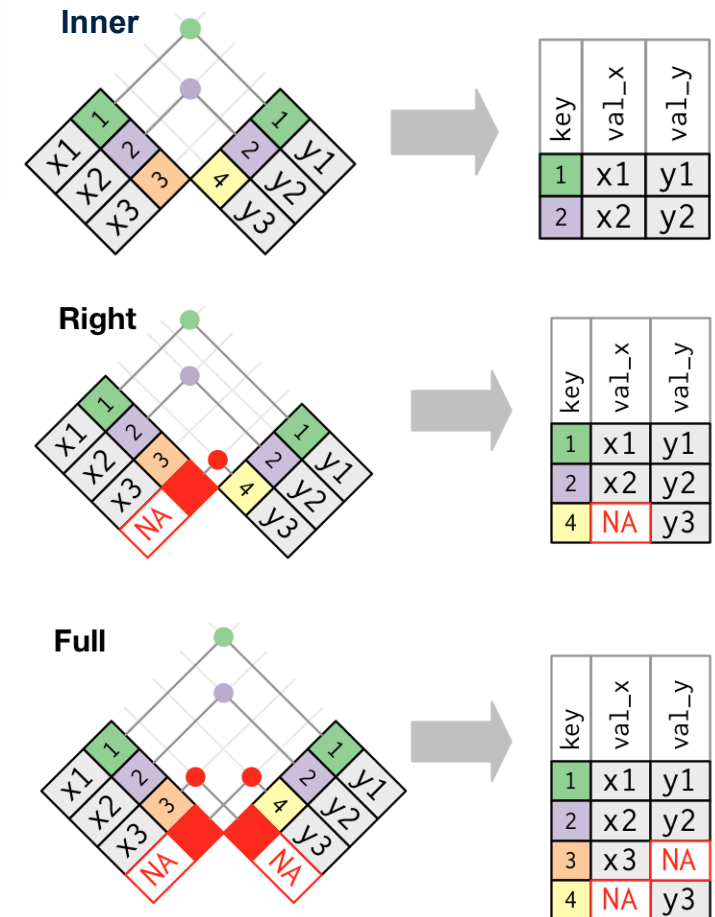
- Mutating joins allows us to combine pairs of tables. Matches observations by their keys and then copies variables across.
- Like *mutate()* – *left\_join()* adds variables to the end. The argument *by* is the key.
- Suppose we wanted to put the full name of the airlines into the flights table.
- We can do this by matching by the carrier column.

```
> flights2 = left_join(select(flights, carrier, dep_time, dep_delay, flight),
+ airlines, by = "carrier")
> print(flights2)
A tibble: 336,776 x 5
 carrier dep_time dep_delay flight name
 <chr> <int> <dbl> <int> <chr>
1 UA 517 2 1545 United Air Lines Inc.
2 UA 533 4 1714 United Air Lines Inc.
3 AA 542 2 1141 American Airlines Inc.
4 B6 544 -1 725 JetBlue Airways
5 DL 554 -6 461 Delta Air Lines Inc.
6 UA 554 -4 1696 United Air Lines Inc.
7 B6 555 -5 507 JetBlue Airways
8 EV 557 -3 5708 ExpressJet Airlines Inc.
9 B6 557 -3 79 JetBlue Airways
10 AA 558 -2 301 American Airlines Inc.
... with 336,766 more rows
```




# Inner & Outer Joins

- *left\_join(x,y)* includes observations with keys in *x*, but not necessarily in *y*.
- There are alternatives:
  - *inner\_join(x,y)* includes observations with keys in both *x* and *y*.
  - *right\_join(x,y)* includes observations with keys in *y*, but not necessarily in *x*.
  - *full\_join(x,y)* keeps all observations with keys in either in *x* or in *y*.




# Keys

- Duplicate keys in one table – *left\_join(x,y)* will keep all duplicate keys in x.



- Duplicate keys in both tables – *left\_join(x,y)* will create all possible combinations.



# Keys (cont.)

- If you ignore the keys argument the left\_join() will match by common variables that exists in both tables.
  - For e.g. the flights weather tables have origin, year, month, day and hour.

```
> left_join(flights, weather)
Joining, by = c("year", "month", "day", "origin", "hour", "time_hour")
A tibble: 336,776 x 28
 year month day dep_time sched_dep_time dep_delay arr_time sched_arr_time arr_delay carrier flight tailnum origin
 <int> <int> <int> <int> <int> <dbl> <int> <int> <dbl> <chr> <int> <chr> <chr>
1 2013 1 1 517 515 2 830 819 11 UA 1545 N14228 EWR
2 2013 1 1 533 529 4 850 830 20 UA 1714 N24211 LGA
3 2013 1 1 542 540 2 923 850 33 AA 1141 N619AA JFK
4 2013 1 1 544 545 -1 1004 1022 -18 B6 725 N804JB JFK
5 2013 1 1 554 600 -6 812 837 -25 DL 461 N668DN LGA
6 2013 1 1 554 558 -4 740 728 12 UA 1696 N39463 EWR
7 2013 1 1 555 600 -5 913 854 19 B6 507 N516JB EWR
8 2013 1 1 557 600 -3 709 723 -14 EV 5708 N829AS LGA
9 2013 1 1 557 600 -3 838 846 -8 B6 79 N593JB JFK
10 2013 1 1 558 600 -2 753 745 8 AA 301 N3ALAA LGA
... with 336,766 more rows, and 15 more variables: dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
minute <dbl>, time_hour <dtm>, temp <dbl>, dewp <dbl>, humid <dbl>, wind_dir <dbl>, wind_speed <dbl>,
wind_gust <dbl>, precip <dbl>, pressure <dbl>, visib <dbl>
```

# Keys (cont..)

If there are multiple keys in table *x* to join with table *y*:

- Specify using *by = c("a" = "b")* where column *a* in table *x* matched to column *b* in table *y*.
- For e.g. if we want to connect flights data to airport data, we need decide whether to use origin or dest to connect to faa.

*left\_join(flights, airports, by = c("dest" = "faa"))* will append destination airport data to the flight table.

*left\_join(flights, airports, by = c("origin" = "faa"))* will append origin airport data to the flight table.

- Assign the expressions to variables and look at the last 7 columns.

# Combining Tables : Filtering Joins

- Filtering joins affect the observations in the tables.
  - `semi_join(x,y)` keeps all observations in table x that matches y.
  - `anti_join(x,y)` removes all observations in table x that matches y.
- For e.g. suppose you wanted to see the data for flights flying to the top 10 destinations.

```
> (top_dest <- head(count(flights, dest, sort = TRUE), 10))
```

```
A tibble: 10 x 2
```

|    | dest  | n     |
|----|-------|-------|
|    | <chr> | <int> |
| 1  | ORD   | 17283 |
| 2  | ATL   | 17215 |
| 3  | LAX   | 16174 |
| 4  | BOS   | 15508 |
| 5  | MCO   | 14082 |
| 6  | CLT   | 14064 |
| 7  | SFO   | 13331 |
| 8  | FLL   | 12055 |
| 9  | MIA   | 11728 |
| 10 | DCA   | 9705  |

```
> semi_join(flights, top_dest, by = "dest")
```

```
A tibble: 141,145 x 19
```

|    | year  | month | day   | dep_time | sched_dep_time | dep_delay | arr_time | sched_arr_time | arr_delay |
|----|-------|-------|-------|----------|----------------|-----------|----------|----------------|-----------|
|    | <int> | <int> | <int> | <int>    | <int>          | <dbl>     | <int>    | <int>          | <dbl>     |
| 1  | 2013  | 1     | 1     | 542      | 540            | 2         | 923      | 850            | 33        |
| 2  | 2013  | 1     | 1     | 554      | 600            | -6        | 812      | 837            | -25       |
| 3  | 2013  | 1     | 1     | 554      | 558            | -4        | 740      | 728            | 12        |
| 4  | 2013  | 1     | 1     | 555      | 600            | -5        | 913      | 854            | 19        |
| 5  | 2013  | 1     | 1     | 557      | 600            | -3        | 838      | 846            | -8        |
| 6  | 2013  | 1     | 1     | 558      | 600            | -2        | 753      | 745            | 8         |
| 7  | 2013  | 1     | 1     | 558      | 600            | -2        | 924      | 917            | 7         |
| 8  | 2013  | 1     | 1     | 558      | 600            | -2        | 923      | 937            | -14       |
| 9  | 2013  | 1     | 1     | 559      | 559            | 0         | 702      | 706            | -4        |
| 10 | 2013  | 1     | 1     | 600      | 600            | 0         | 851      | 858            | -7        |

# ... with 141,135 more rows, and 10 more variables: carrier <chr>, flight <int>, tailnum <chr>,  
# origin <chr>, dest <chr>, air\_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,  
# time\_hour <dtm>

# Combining Tables : Filtering Joins (cont.)

- On the other hand, if you wanted to see the data for flights NOT flying to the top 10 destinations.
- Useful to find mismatches. For e.g. connecting flights and planes tables - find observations in flights table that don't exist in planes using "tailnum".

```
> anti_join(flights, top_dest, by = "dest")
A tibble: 195,631 x 19
 year month day dep_time sched_dep_time dep_delay arr_time sched_arr_time arr_delay
 <int> <int> <int> <int> <int> <dbl> <int> <int> <dbl>
1 2013 1 1 517 515 2 830 819 11
2 2013 1 1 533 529 4 850 830 20
3 2013 1 1 544 545 -1 1004 1022 -18
4 2013 1 1 557 600 -3 709 723 -14
5 2013 1 1 558 600 -2 849 851 -2
6 2013 1 1 558 600 -2 853 856 -3
7 2013 1 1 559 600 -1 941 910 31
8 2013 1 1 559 600 -1 854 902 -8
9 2013 1 1 601 600 1 844 850 -6
10 2013 1 1 602 610 -8 812 820 -8
... with 195,621 more rows, and 10 more variables: carrier <chr>, flight <int>, tailnum <chr>,
origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,
> anti_join(flights, planes, by = "tailnum")
A tibble: 52,606 x 19
 year month day dep_time sched_dep_time dep_delay arr_time sched_arr_time arr_delay carrier flight
 <int> <int> <int> <int> <int> <dbl> <int> <int> <dbl> <chr> <int>
1 2013 1 1 558 600 -2 753 745 8 AA 301
2 2013 1 1 559 600 -1 941 910 31 AA 707
3 2013 1 1 600 600 0 837 825 12 MQ 4650
4 2013 1 1 602 605 -3 821 805 16 MQ 4401
5 2013 1 1 608 600 8 807 735 32 MQ 3768
6 2013 1 1 611 600 11 945 931 14 UA 303
7 2013 1 1 623 610 13 920 915 5 AA 1837
8 2013 1 1 624 630 -6 840 830 10 MQ 4599
9 2013 1 1 628 630 -2 1137 1140 -3 AA 413
10 2013 1 1 629 630 -1 824 810 14 AA 303
... with 52,596 more rows, and 8 more variables: tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>,
distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

# Set Operations without dplyr

- There 4 set operations in base R that could be used as well:
  - *intersect(x,y)* returns observations that are in both x and y.
  - *union(x,y)* returns unique observations in x and y.
  - *setdiff(x,y)* returns observations in x but not in y.
  - *merge(x,y)* same as mutate join functions –
    - *merge(x,y)*                       $\hookrightarrow$                       *inner\_join(x,y)*
    - *merge(x,y, all.x = TRUE)*                       $\hookrightarrow$                       *left\_join(x,y)*
    - *merge(x,y, all.y = TRUE)*                       $\hookrightarrow$                       *right\_join(x,y)*
    - *merge(x,y, , all.x = TRUE, all.y = TRUE)*                       $\hookrightarrow$                       *full\_join(x,y)*.

# Example set 1

- Load “hmda\_2017\_tx\_all\_40.csv” and “hmda\_2017\_tx\_all\_06.csv”.
- Sample 500 observations from *hmda\_2017\_tx\_all\_06*.
- Add the lien status from *hmda\_2017\_tx\_all\_40* to the sample from *\_06*.
- Which lien status is most common?



# Example set 2

Install package *Lahman* and load it. There are many data frames (df).

See *?Managers* and *?AwardsManagers*.

1. Attach the *awardID* from *AwardsManagers* to the *Managers* df.

See *?Salaries* and *?AwardsPlayers*.

2. Attach the *awardID* from *AwardsPlayers* to *Salaries* df.
3. See the top and bottom 10 salaries of players and compare awards.
4. See *?Appearances*. Are there any players who received salary in a year but never appeared for a game?
5. See *?People*. Are there any players who did not show up on salaries data, yet played games? If so, how many games? Arrange in descending order.