



IS 6733: Deep Learning on Cloud Platforms

01 Artificial Neural Networks

Copy Right Notice






- **Most slides in this presentation are adopted from slides of text book and various sources. The Copyright belong to the original authors. Thanks!**

Why Neural Network

- 🐾 **Some tasks can be done easily by humans but are hard by conventional paradigms on Von Neumann machine with algorithmic approach**
 - 🐾 Pattern recognition (old friends, hand-written characters)
 - 🐾 Content addressable recall
 - 🐾 Approximate, common sense reasoning (driving, playing piano, baseball player)
- 🐾 **These tasks are often experience based, hard to apply logic.**

Biological Motivation

Humans:

-  Neuron switching time ~ 0.001 second
-  Number of neurons $\sim 10^{10}$
-  Connections per neuron $\sim 10^{4-5}$
-  Scene recognition time ~ 0.1 second
-  Highly parallel computation process.

 **Biological Learning Systems are built of very complex webs of interconnected neurons.**

 **Information-Processing abilities of biological neural systems must follow from highly parallel processes operating on representations that are distributed over many neurons**

What is an neural network

- 🐾 **A set of nodes (units, neurons, processing elements)**
 - 🐾 Each node has input and output
 - 🐾 Each node performs a simple computation by its **node function**
- 🐾 **Weighted connections between nodes**
 - 🐾 Connectivity gives the structure/architecture of the net
 - 🐾 What can be computed by a NN is primarily determined by the connections and their weights
- 🐾 **A very much simplified version of networks of neurons in animal nerve systems**

ANN vs. Bio NN

ANN

- **Nodes**
 - input
 - output
 - node function
- **Connections**
 - connection strength

Bio NN

- **Cell body**
 - signal from other neurons
 - firing frequency
 - firing mechanism
- **Synapses**
 - synaptic strength

Properties of artificial neural nets

- 🐾 **Many neuron-like threshold switching units**
- 🐾 **Many weighted interconnections among units**
- 🐾 **Highly parallel, distributed process**
- 🐾 **Emphasis on tuning weights automatically**

When to Consider Neural Networks

- 🐾 **Input is high-dimensional discrete or real-valued**
- 🐾 **Output is discrete or real valued**
- 🐾 **Output is a vector of values**
- 🐾 **Possibly noisy data**
- 🐾 **Form of target function is unknown**
- 🐾 **Human readability of result is unimportant**
- 🐾 **Examples:**
 - 🐾 Speech phoneme recognition
 - 🐾 Image classification
 - 🐾 Financial prediction

History of Neural Networks

- 🐾 **1943: McCulloch and Pitts proposed a model of a neuron --> Perceptron**
- 🐾 **1960s: Widrow and Hoff explored Perceptron networks (which they called “Adelines”) and the delta rule.**
- 🐾 **1962: Rosenblatt proved the convergence of the perceptron training rule.**
- 🐾 **1969: Minsky and Papert showed that the Perceptron cannot deal with nonlinearly-separable data sets---even those that represent simple function such as X-OR.**
- 🐾 **1970-1985: Very little research on Neural Nets**
- 🐾 **1986: Invention of Backpropagation [Rumelhart and McClelland, but also Parker and earlier on: Werbos] which can learn from nonlinearly-separable data sets.**
- 🐾 **Since 1985: A lot of research in Neural Nets!**

A Perceptron (a neuron)

🌸 The network

- 🌸 Input vector \mathbf{i}_j (including threshold input = 1)
- 🌸 Weight vector $\mathbf{w} = (w_0, w_1, \dots, w_n)$

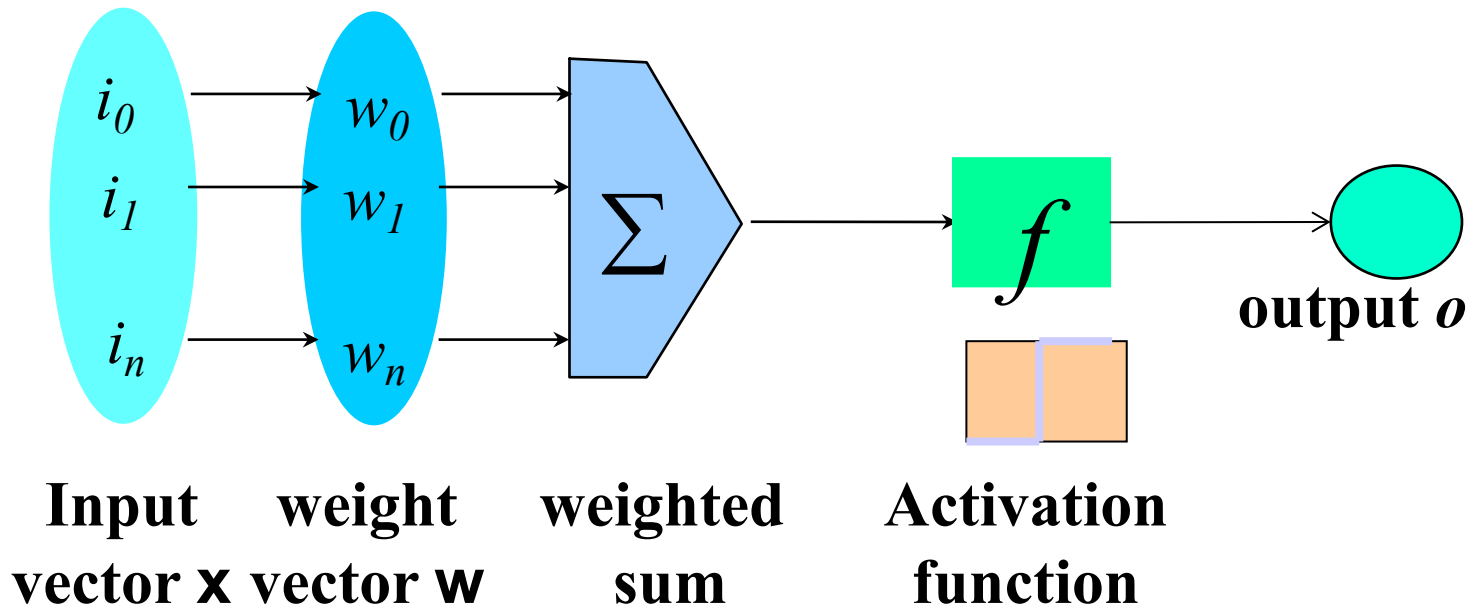
$$net = \mathbf{w} \cdot \mathbf{i}_j = \sum_{k=0}^n w_k i_{k,j}$$

- 🌸 Output: bipolar (-1, 1) using the sign node function



$$output = \begin{cases} 1 & \text{if } \mathbf{w} \cdot \mathbf{i}_j > 0 \\ -1 & \text{otherwise} \end{cases}$$

🌸 Training samples

- 🌸 Pairs $(\mathbf{i}_j, class(\mathbf{i}_j))$ where $class(\mathbf{i}_j)$ is the correct classification of \mathbf{i}_j




Aside: Multilayer Perceptron



Activation functions


🐾 Step (threshold) function

$$f(\text{net}) = \begin{cases} a & \text{if } \text{net} < c \\ b & \text{if } \text{net} > c \end{cases}$$



🐾 Ramp function

$$f(\text{net}) = \begin{cases} a & \text{if } \text{net} \leq c \\ b & \text{if } \text{net} \geq d \\ a + \frac{(\text{net} - c)(b - a)}{(d - c)} & \text{otherwise} \end{cases}$$




Activation functions

🐾 Sigmoid function

- 🐾 S-shaped
- 🐾 Continuous and everywhere differentiable
- 🐾 Rotationally symmetric about some point ($net = c$)
- 🐾 Asymptotically approaches saturation points

$$f(net) = z + \frac{1}{1 + \exp(-x \cdot net + y)}$$

$$f(net) = \tanh(x \cdot net - y) + z,$$




Decision Surface of a Perceptron: Linear separability

🐾 **n dimensional patterns (x_1, \dots, x_n)**

🐾 Hyperplane $w_0 + w_1 x_1 + w_2 x_2 + \dots + w_n x_n = 0$ dividing the space into two regions

🐾 **Can we get the weights from a set of sample patterns?**

🐾 If the problem is linearly separable, then YES (by perceptron learning)




Examples of linearly separable classes



Logical AND function

patterns (bipolar) decision boundary

x1	x2	output	w1 = 1
-1	-1	-1	w2 = 1
-1	1	-1	w0 = -1
1	-1	-1	
1	1	1	$-1 + x1 + x2 = 0$




x: class I (output = 1)
o: class II (output = -1)



Logical OR function

patterns (bipolar) decision boundary

x1	x2	output	w1 = 1
-1	-1	-1	w2 = 1
-1	1	1	w0 = 1
1	-1	1	
1	1	1	$1 + x1 + x2 = 0$




x: class I (output = 1)
o: class II (output = -1)

Functions not representable

🐾 **Some functions are not representable by perceptron**

🐾 Not linearly separable



Perceptron Training Rule

🐾 Training:

- 🐾 Update \mathbf{w} so that all sample inputs are correctly classified (if possible)
- 🐾 If an input \mathbf{i}_j is misclassified by the current \mathbf{w}
 - 🐾 $\text{class}(\mathbf{i}_j) \cdot \mathbf{w} \cdot \mathbf{i}_j < 0$
 - 🐾 change \mathbf{w} to $\mathbf{w} + \Delta\mathbf{w}$ so that $(\mathbf{w} + \Delta\mathbf{w}) \cdot \mathbf{i}_j$ is closer to $\text{class}(\mathbf{i}_j)$

🐾 Perceptron Training Rule

$$w_i = w_i + \Delta w_i$$

🐾 Where

$$\Delta w_i = \eta(t - o)x_i$$

🐾 Where

- 🐾 $t = c(\vec{x})$ is the target value
- 🐾 o is perceptron output
- 🐾 η is a small positive constant, called learning rate

Perceptron Training Algorithm

- 🐾 **Start with a randomly chosen weight vector w_0**
- 🐾 **Let $k=1$;**
- 🐾 **While some input vectors remain misclassified , do**
 - 🐾 Let x_j be a misclassified input vector
 - 🐾 Update the weight vector to $w_k = w_{k-1} + \eta(t - o)x_k$
 - 🐾 Increment k ;
- 🐾 **End while**

Perceptron Training Rule

🐾 It will converge if

🐾 Training data is linearly separable

🐾 η is a sufficiently small

🐾 **Theorem:** If there is a w^* such that $f(i_p \cdot w^*) = class(i_p)$ for all P training sample patterns $\{i_p, class(i_p)\}$, then for any start weight vector w^0 , the perceptron learning rule will converge to a weight vector w^+ such that for all p

$$f(i_p \cdot w^+) = class(i_p)$$

(w^* and w^+ may not be the same.)

Perceptron Training Rule

Justification

$$(w + \eta \cdot (t - o) \cdot x_k) \cdot x_k = w \cdot x_k + \eta \cdot (t - o) \cdot x_k \cdot x_k$$

then

$$(w + \eta \cdot (t - o) \cdot x_k) \cdot x_k - w \cdot x_k = \eta \cdot (t - o) \cdot x_k \cdot x_k$$

since $x_k \cdot x_k > 0$

$$\begin{cases} > 0 & \text{if } \text{class}(i_j) = 1 \\ < 0 & \text{if } \text{class}(i_j) = -1 \end{cases}$$

\Rightarrow new *net* moves toward $\text{class}(i_j)$

Perceptron Training Rule


 **Termination criteria: learning stops when all samples are correctly classified**

 Assuming the problem is linearly separable

 Assuming the learning rate (η) is sufficiently small

 **Choice of learning rate:**

 If η is too large: existing weights are overtaken by Δw

 If η is too small (≈ 0): very slow to converge

 Common choice: $0.1 < \eta < 1$.

Example, perceptron learning function AND

Training samples

	in_0	in_1	in_2	d
p0	1	-1	-1	-1
p1	1	-1	1	-1
p2	1	1	-1	-1
p3	1	1	1	1

Initial weights $W(0)$

w0	w1	w2
1	1	-1

Learning rate = 1

- **Present p0**
 - $\text{net} = W(0)p0 = (1, 1, -1)(1, -1, -1) = 1$
 - p0 misclassified, learning occurs
 - $W(1) = W(0) + (t-o)*p0 = (-1, 3, 1)$
 - New net = $W(1)p0 = -5$ is closer to target ($t = -1$)
- **Present p1**
 - $\text{net} = (-1, 3, 1)(1, -1, 1) = -3$
 - no learning occurs
- **Present p2**
 - $\text{net} = (-1, 3, 1)(1, 1, -1) = 1$
 - $W(2) = (-1, 3, 1) + (-2)(1, 1, -1) = (-3, 1, 3)$
 - New net = $W(2)p2 = -5$
- **Present p3**
 - $\text{net} = (-3, 1, 3)(1, 1, 1) = 1$
 - no learning occurs
- **Present p0, p1, p2, p3**
 - All correctly classified with $W(2)$
 - Learning stops with $W(2)$

Delta Rule

- 🐾 **The perceptron rule fail to converge if the examples are not linearly separable.**
- 🐾 **Delta rule will converge toward a best-fit approximation to the target concept if the training example are not linearly separable.**
 - 🐾 **The delta rule is to use gradient descent to search the hypothesis space.**

Gradient Descent

🐾 Consider simpler linear unit, where

$$o(x) = \vec{w} \cdot \vec{x} = w_0 + w_1 x_1 + w_2 x_2 + \dots + w_n x_n$$

🐾 Let's learn w_i 's that minimize the squared error

$$E(\vec{w}) = \frac{1}{2} \sum_{d \in D} (t_d - o_d)^2$$

🐾 Where D is the set of training examples.

Gradient Descent

Gradient


$$\nabla E[\vec{w}] \equiv \left[\frac{\partial E}{\partial w_0}, \frac{\partial E}{\partial w_1}, \dots, \frac{\partial E}{\partial w_n} \right]$$

Training rule:

$$\Delta \vec{w} = -\eta \nabla E[\vec{w}]$$

 i.e.,

$$\Delta w_i = -\eta \frac{\partial E}{\partial w_i}$$



Gradient Descent

$$\begin{aligned}\frac{\partial E}{\partial w_i} &= \frac{\partial}{\partial w_i} \frac{1}{2} \sum_d (t_d - o_d)^2 \\ &= \frac{1}{2} \sum_d \frac{\partial}{\partial w_i} (t_d - o_d)^2 \\ &= \frac{1}{2} \sum_d 2(t_d - o_d) \frac{\partial}{\partial w_i} (t_d - o_d) \\ &= \sum_d (t_d - o_d) \frac{\partial}{\partial w_i} (t_d - \vec{w} \cdot \vec{x}_d) \\ \frac{\partial E}{\partial w_i} &= \sum_d (t_d - o_d) (-x_{i,d})\end{aligned}$$

$$\Delta w_i = \eta \sum_{d \in D} (t_d - o_d) x_{id}$$

Gradient Descent

GRADIENT-DESCENT(*training_examples*, η)

Each training example is a pair of the form $\langle \vec{x}, t \rangle$, where \vec{x} is the vector of input values, and t is the target output value. η is the learning rate (e.g., .05).

- Initialize each w_i to some small random value
- Until the termination condition is met, Do
 - Initialize each Δw_i to zero.
 - For each $\langle \vec{x}, t \rangle$ in *training_examples*, Do
 - * Input the instance \vec{x} to the unit and compute the output o
 - * For each linear unit weight w_i , Do



$$\Delta w_i \leftarrow \Delta w_i + \eta(t - o)x_i$$

- For each linear unit weight w_i , Do






$$w_i \leftarrow w_i + \Delta w_i$$

Stochastic gradient descent

Practical difficulties of gradient descent





-  Converge to local minimum can sometimes be quite slow
-  If there are multiple local minima in the error surface, there is no guarantee that the procedure will find the global minimum.

Stochastic gradient descent: update weights incrementally

-  Do until satisfied
 -  For each training example d in D
 -  Compute the gradient $\nabla E_d[\vec{x}]$
 -  Then, $\vec{w} = \vec{w} - \eta \nabla E_d[\vec{w}]$
-  Stochastic (incremental) gradient descent can approximate standard gradient descent arbitrarily closely if learning rate made small enough.

Stochastic gradient descent

Key differences:

-  In standard gradient descent, the error is summed over all examples before updating weights, where in stochastic gradient weights are updated upon examining each training example
-  Summing over multiple examples in standard gradient descent requires more computation per weight update step
 -  Use larger step size per weight in standard gradient descent
-  In cases where there are multiple local minima with respect to $E(w)$, stochastic gradient descent can sometimes avoid falling into these local minima.

Summary

- 🐾 **Perceptron training rule updates weights on the error in the thresholded perceptron output**



$$o(\vec{x}) = \text{sgn}(\vec{w} \cdot \vec{x})$$

- 🐾 **Delta training rule updates weights on the error in the unthresholded linear combination of inputs**





$$o(\vec{x}) = \vec{w} \cdot \vec{x}$$

Summary


Perceptron training rule guaranteed to succeed if

-  Training examples are linearly separable
-  Sufficiently small learning rate

Delta training rule uses gradient descent

-  Guaranteed to converge to hypothesis with minimum squared error
-  Given sufficiently small learning rate
-  Even when training data contains noise
-  Even when training data not separable by H.

A Multilayer Neural Network



How A Multilayer Neural Network Works?

- 🐾 The inputs to the network correspond to the attributes measured for each training example
- 🐾 Inputs are fed simultaneously into the units making up the input layer
- 🐾 They are then weighted and fed simultaneously to a hidden layer
- 🐾 The number of hidden layers is arbitrary, although usually only one
- 🐾 The weighted outputs of the last hidden layer are input to units making up the output layer, which emits the network's prediction
- 🐾 The network is feed-forward in that none of the weights cycles back to an input unit or to an output unit of a previous layer
- 🐾 From a statistical point of view, networks perform **nonlinear regression**: Given enough hidden units and enough training samples, they can closely approximate any function

Multilayer Networks of Sigmoid Units

🐾 Architecture:

🐾 **Feedforward** network of at least one layer of **non-linear** hidden nodes, e.g., # of layers $L \geq 2$ (not counting the input layer)

🐾 Node function is differentiable

🐾 most common: **sigmoid function**

$$S(net) = \frac{1}{1 + e^{(-net)}}$$

🐾 **Nice property:**

$$\frac{dS(x)}{dx} = S(x)(1 - S(x))$$









🐾 **We can derive gradient descent rules to train**

🐾 **One sigmoid unit**

🐾 **Multilayer networks of sigmoid units**

Backpropagation Learning

Notation:

-  x_{ji} : the i th input to unit j
-  w_{ji} : the weight associated with i th input to unit j
-  $\text{net}_j = \sum_i w_{ji} x_{ji}$ (the weighted sum of inputs for unit j)
-  o_j : the output computed by unit j
-  t_j : the target output for unit j
-  σ : the sigmoid function
-  outputs: the set of units in the final layer of the network
-  $\text{Downstream}(j)$: the set of units whose immediate inputs include the output of unit j .

Backpropagation Learning

🐾 Idea of BP learning:

- 🐾 Update of weights in w_{21} (from hidden layer to output layer): delta rule as in a single layer net using sum square error
- 🐾 Delta rule is not applicable to updating weights in w_{10} (from input and hidden layer) because we don't know the desired values for hidden nodes
- 🐾 **Solution:** Propagating errors at output nodes down to hidden nodes, these computed errors on hidden nodes drives the update of weights in w_{10} (again by delta rule), thus called error **BACKPROPAGATION (BP)** learning
- 🐾 How to compute errors on hidden nodes is the key
- 🐾 Error backpropagation can be continued downward if the net has more than one hidden layer
- 🐾 Proposed first by Werbos (1974), current formulation by Rumelhart, Hinton, and Williams (1986)

Backpropagation Learning

- 🐾 For each training example d every weight w_{ji} is updated by adding to it Δw_{ji}

$$\Delta w_{ji} = -\eta \frac{\partial E_d}{\partial w_{ji}}$$

- 🐾 Where E_d is the error on training example d , summed over all output units in the network

$$E(\vec{w}) = \frac{1}{2} \sum_{d \in D} (t_d - o_d)^2$$

Backpropagation Learning

- 🐾 Noted that weight w_{ji} can influence the rest of the network only through net_j . Therefore, we can use the **chain rule** to write

$$\frac{\partial E_d}{\partial w_{ji}} = \frac{\partial E_d}{\partial net_j} \frac{\partial net_j}{\partial w_{ji}} = \frac{\partial E_d}{\partial net_j} x_{ji}$$

- 🐾 Our remaining task is to derive a convenient expression of $\frac{\partial E_d}{\partial net_j}$. Two cases are considered:
 - 🐾 Unit j is an output unit for the network
 - 🐾 Unit j is an internal unit.

Backpropagation Learning

Training rule for output unit weights

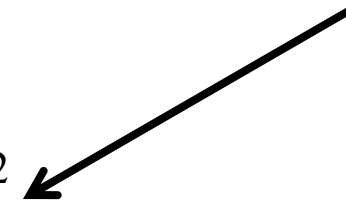
 net_j can influence the rest of the network only through o_j ,
Then

$$\frac{\partial E_d}{\partial net_j} = \frac{\partial E_d}{\partial o_j} \frac{\partial o_j}{\partial net_j}$$

 First term:

$$\begin{aligned} \frac{\partial E_d}{\partial o_j} &= \frac{\partial}{\partial o_j} \frac{1}{2} \sum_{k \in \text{outputs}} (t_k - o_k)^2 \\ &= \frac{\partial}{\partial o_j} \frac{1}{2} (t_j - o_j)^2 = \frac{1}{2} \times 2 \times (t_j - o_j) \frac{\partial (t_j - o_j)}{\partial o_j} \\ &= -(t_j - o_j) \end{aligned}$$

Derivatives will
be zero for all
output units
except j



Backpropagation Learning

🐾 Second term:

$$\frac{\partial o_j}{\partial net_j} = \frac{\partial \sigma(net_j)}{\partial net_j} = o_j(1 - o_j)$$

$o_j = \sigma(net_j)$ ←

🐾 Put it together:


$$\frac{\partial E_d}{\partial net_j} = -(t_j - o_j)o_j(1 - o_j)$$

🐾 Then, we have the stochastic gradient descent rule for output units

$$\Delta w_{ji} = -\eta \frac{\partial E_d}{\partial w_{ji}} = \eta(t_j - o_j)o_j(1 - o_j)x_{ji}$$

Backpropagation Learning

Training rule for hidden unit weights

 net_j (j is the internal node) can influence the rest of the network through $Downstream(j)$, Then

$$\begin{aligned}\frac{\partial E_d}{\partial net_j} &= \sum_{k \in Downstream(j)} \frac{\partial E_d}{\partial net_k} \frac{\partial net_k}{\partial net_j} \\ &= \sum_{k \in Downstream(j)} -\delta_k \frac{\partial net_k}{\partial net_j} \\ &= \sum_{k \in Downstream(j)} -\delta_k \frac{\partial net_k}{\partial o_j} \frac{\partial o_j}{\partial net_j} \\ &= \sum_{k \in Downstream(j)} -\delta_k w_{kj} \frac{\partial o_j}{\partial net_j} \\ &= \sum_{k \in Downstream(j)} -\delta_k w_{kj} o_j (1 - o_j)\end{aligned}$$

Backpropagation Learning

 **We set**

$$\delta_j = -\frac{\partial E_d}{\partial \text{net}_j} = o_j(1 - o_j) \sum_{k \in \text{Downstream}(j)} \delta_k w_{kj}$$

 **Then, we have the stochastic gradient descent rule for hidden units**

$$\Delta w_{ji} = \eta \delta_j x_{ji}$$

Backpropagation Learning

BACKPROPAGATION(*training_examples*, η , n_{in} , n_{out} , n_{hidden})

Each training example is a pair of the form $\langle \vec{x}, \vec{t} \rangle$, where \vec{x} is the vector of network input values, and \vec{t} is the vector of target network output values.

η is the learning rate (e.g., .05). n_{in} is the number of network inputs, n_{hidden} the number of units in the hidden layer, and n_{out} the number of output units.

The input from unit i into unit j is denoted x_{ji} , and the weight from unit i to unit j is denoted w_{ji} .

- Create a feed-forward network with n_{in} inputs, n_{hidden} hidden units, and n_{out} output units.
- Initialize all network weights to small random numbers (e.g., between $-.05$ and $.05$).
- Until the termination condition is met, Do
 - For each $\langle \vec{x}, \vec{t} \rangle$ in *training_examples*, Do

Propagate the input forward through the network:

1. Input the instance \vec{x} to the network and compute the output o_u of every unit u in the network.

Propagate the errors backward through the network:

2. For each network output unit k , calculate its error term δ_k

$$\delta_k \leftarrow o_k(1 - o_k)(t_k - o_k) \quad (\text{T4.3})$$

3. For each hidden unit h , calculate its error term δ_h

$$\delta_h \leftarrow o_h(1 - o_h) \sum_{k \in \text{outputs}} w_{kh} \delta_k \quad (\text{T4.4})$$

4. Update each network weight w_{ji}

$$w_{ji} \leftarrow w_{ji} + \Delta w_{ji}$$

where

$$\Delta w_{ji} = \eta \delta_j x_{ji} \quad (\text{T4.5})$$


Learning Hidden Layer Representations

🐾 A target function

Input	Output
10000000	→ 10000000
01000000	→ 01000000
00100000	→ 00100000
00010000	→ 00010000
00001000	→ 00001000
00000100	→ 00000100
00000010	→ 00000010
00000001	→ 00000001


Learning Hidden Layer Representations

🐾 A network:




Learning Hidden Layer Representations

🐾 Sum of squared errors for each output unit




Learning Hidden Layer Representations

🐾 Hidden unit encoding for input 01000000



Learning Hidden Layer Representations

🐾 **Weights from inputs to on hidden unit**



Learning Hidden Layer Representations

🐾 **Learned hidden layer representation after 5000 training epochs**


Input		Hidden Values		Output
100000000	→	.89 .04 .08	→	100000000
010000000	→	.01 .11 .88	→	010000000
001000000	→	.01 .97 .27	→	001000000
000100000	→	.99 .97 .71	→	000100000
000010000	→	.03 .05 .02	→	000010000
000001000	→	.22 .99 .99	→	000001000
000000100	→	.80 .01 .98	→	000000100
000000001	→	.60 .94 .01	→	000000001

Example, BP learning function XOR

Training samples (bipolar)

	in_1	in_2	d
P0	-1	-1	-1
P1	-1	1	1
P2	1	-1	1
P3	1	1	-1

Network: 2-2-1 with thresholds (fixed output 1)



- Initial weights $W(0)$

$$w_1^{(1,0)} : (-0.5, 0.5, -0.5)$$

$$w_2^{(1,0)} : (-0.5, -0.5, 0.5)$$

$$w^{(2,1)} : (-1, 1, 1)$$

- Learning rate = 0.2
- Node function: hyperbolic tangent

$$g(x) = \tanh(x) = \frac{1 - e^{-x}}{1 + e^{-x}};$$

$$\lim_{x \rightarrow \pm\infty} g(x) = \pm 1$$

$$s(x) = \frac{1}{1 + e^{-x}};$$

$$g(x) = 2s(x) - 1$$

$$s'(x) = s(x)(1 - s(x))$$

$$g'(x) = 0.5(1 + g(x))(1 - g(x))$$

Present $\mathbf{P}_0 = (1, -1, -1)$: $\mathbf{d}_0 = -1$

Forward computing

$$net_1 = w_1^{(1,0)} p_0 = (-0.5, 0.5, -0.5) (1, -1, -1) = -0.5$$

$$net_2 = w_2^{(1,0)} p_0 = (-0.5, -0.5, 0.5) (1, -1, -1) = -0.5$$

$$x_1^{(1)} = g(net_1) = 2/(1 + e^{0.5}) - 1 = -0.24492$$

$$x_2^{(1)} = g(net_2) = 2/(1 + e^{0.5}) - 1 = -0.24492$$

$$net_o = w^{(2,1)} x^{(1)} = (-1, 1, 1)(1, -0.24492, -0.24492) = -1.48984$$

$$o = g(net_o) = -0.63211$$

Error back propogating

$$l = d - o = -1 - (-0.63211) = -0.36789$$

$$\begin{aligned} \delta &= l \cdot g'(net_o) = l \cdot (1 + g(net_o))(1 - g(net_o)) \\ &= -0.3679 \cdot (1 - 0.6321)(1 + 0.6321) = -0.2209 \end{aligned}$$

$$\begin{aligned} \mu_1 &= \delta \cdot w_1^{(2,1)} \cdot g'(net_1) \\ &= -0.2209 \cdot 1 \cdot (1 - 0.24492) \cdot (1 + 0.24492) = -0.20765 \end{aligned}$$

$$\begin{aligned} \mu_2 &= \delta \cdot w_2^{(2,1)} \cdot g'(net_2) \\ &= -0.2209 \cdot 1 \cdot (1 - 0.24492) \cdot (1 + 0.24492) = -0.20765 \end{aligned}$$

Weight update

$$\Delta w^{(2,1)} = \eta \cdot \delta \cdot x^{(1)}$$

$$= 0.2 \cdot (-0.2209) \cdot (1, -0.2449, -0.2449) = (-0.0442, 0.0108, 0.0108)$$

$$\begin{aligned} w^{(2,1)} &= w^{(2,1)} + \Delta w^{(2,1)} = (-1, 1, 1) + (-0.0442, 0.0108, 0.0108) \\ &= (-0.5415, 1.0108, 1.0108) \end{aligned}$$

$$\Delta w_1^{(1,0)} = \eta \cdot \mu_1 \cdot p_0 = 0.2 \cdot (-0.2077) \cdot (1, -1, -1) = (-0.0415, 0.0415, 0.0415)$$


$$\Delta w_2^{(1,0)} = \eta \cdot \mu_2 \cdot p_0 = 0.2 \cdot (-0.2077) \cdot (1, -1, -1) = (-0.0415, 0.0415, 0.0415)$$

$$\begin{aligned} w_1^{(1,0)} &= w_1^{(1,0)} + \Delta w_1^{(1,0)} = (-0.5, 0.5, -0.5) + (-0.0415, 0.0415, 0.0415) \\ &= (-0.5415, 0.5415, -0.4585) \end{aligned}$$

$$\begin{aligned} w_2^{(1,0)} &= w_2^{(1,0)} + \Delta w_2^{(1,0)} = (-0.5, -0.5, 0.5) + (-0.0415, 0.0415, 0.0415) \\ &= (-0.5415, -0.4585, 0.5415) \end{aligned}$$

Error for $P_0 = l^2$ reduced from 0.135345 to 0.102823

MSE reduction:
every 10 epochs



Output: every 10 epochs

epoch	1	10	20	40	90	140	190	d
P0	-0.63	-0.05	-0.38	-0.77	-0.89	-0.92	-0.93	-1
P1	-0.63	-0.08	0.23	0.68	0.85	0.89	0.90	1
P2	-0.62	-0.16	0.15	0.68	0.85	0.89	0.90	1
p3	-0.38	0.03	-0.37	-0.77	-0.89	-0.92	-0.93	-1
MSE	1.44	1.12	0.52	0.074	0.019	0.010	0.007	

After epoch 1

	$w_1^{(1,0)}$	$w_2^{(1,0)}$	$w^{(2,1)}$
init	(-0.5, 0.5, -0.5)	(-0.5, -0.5, 0.5)	(-1, 1, 1)
p0	-0.5415, 0.5415, -0.4585	-0.5415, -0.45845, 0.5415	-1.0442, 1.0108, 1.0108
p1	-0.5732, 0.5732, -0.4266	-0.5732, -0.4268, 0.5732	-1.0787, 1.0213, 1.0213
p2	-0.3858, 0.7607, -0.6142	-0.4617, -0.3152, 0.4617	-0.8867, 1.0616, 0.8952
p3	-0.4591, 0.6874, -0.6875	-0.5228, -0.3763, 0.4005	-0.9567, 1.0699, 0.9061



epoch

13	-1.4018, 1.4177, -1.6290	-1.5219, -1.8368, 1.6367	0.6917, 1.1440, 1.1693
40	-2.2827, 2.5563, -2.5987	-2.3627, -2.6817, 2.6417	1.9870, 2.4841, 2.4580
90	-2.6416, 2.9562, -2.9679	-2.7002, -3.0275, 3.0159	2.7061, 3.1776, 3.1667
190	-2.8594, 3.18739, -3.1921	-2.9080, -3.2403, 3.2356	3.1995, 3.6531, 3.6468



Strength of BP

Great representation power




Boolean functions

-  Every Boolean function can be represented by network with single hidden layer
-  But might require exponential hidden units.

Continuous functions

-  Every bounded continuous function can be approximated with arbitrarily small error by network with one hidden layer
-  Any function can be approximated to arbitrary accuracy by a network with two hidden layers

Wide applicability of BP learning

-  Only requires that a good set of training samples is available
-  Does not require substantial prior knowledge or deep understanding of the domain itself (ill structured problems)
-  Tolerates noise and missing data in training samples (graceful degrading)

Easy to implement the core of the learning algorithm

Good generalization power

-  Often produce accurate results for inputs outside the training set

Deficiencies of BP

❀ Learning often takes a long time to converge

- ❀ Complex functions often need hundreds or thousands of epochs

❀ The net is essentially a black box

- ❀ It may provide a desired mapping between input and output vectors (\mathbf{x} , \mathbf{o}) but does not have the information of why a particular \mathbf{x} is mapped to a particular \mathbf{o} .
- ❀ It thus cannot provide an intuitive (e.g., causal) explanation for the computed result.
- ❀ This is because the hidden nodes and the learned weights do not have clear semantics.
 - ❀ What can be learned are operational parameters, not general, abstract knowledge of a domain
- ❀ Unlike many statistical methods, there is no theoretically well-founded way to **assess the quality** of BP learning
 - ❀ What is the confidence level one can have for a trained BP net, with the final E (which may or may not be close to zero)?
 - ❀ What is the confidence level of \mathbf{o} computed from input \mathbf{x} using such net?

Deficiencies of BP

🐾 Problem with gradient descent approach

- 🐾 only guarantees to reduce the total error to a **local minimum**. (E may not be reduced to zero)
- 🐾 Cannot escape from the local minimum error state
- 🐾 **Not every function that is representable can be learned**
- 🐾 How bad: depends on the shape of the error surface. Too many valleys/wells will make it easy to be trapped in local minima
- 🐾 Possible remedies:
 - 🐾 Try nets with different # of hidden layers and hidden nodes (they may lead to different error surfaces, some might be better than others)
 - 🐾 Try different initial weights (different starting points on the surface)
 - 🐾 Forced escape from local minima by random perturbation (e.g., simulated annealing)

Variations of BP nets

🐾 Adding momentum term (to speedup learning)

- 🐾 Weights update at time n contains the momentum of the previous updates, e.g.,

$$\Delta w_{ji}(n) = \eta \delta_j x_{ji} + \alpha \Delta w_{ji}(n-1)$$

- 🐾 Avoid sudden change of directions of weight update (smoothing the learning process)
- 🐾 Error is no longer monotonically decreasing

🐾 Batch mode of weight update


- 🐾 Weight update once per each epoch (cumulated over all P samples)
- 🐾 Smoothing the training sample outliers
- 🐾 Learning independent of the order of sample

Variations of BP nets


🐾 Variations on learning rate η

- 🐾 Fixed rate much smaller than 1
- 🐾 Start with large η , gradually decrease its value
- 🐾 Start with a small η , steadily double it until MSE start to increase
- 🐾 Give known underrepresented samples higher rates
- 🐾 Find the maximum safe step size at each stage of learning (to avoid overshoot the minimum E when increasing η)
- 🐾 **Adaptive learning rate** (delta-bar-delta method)
 - 🐾 Each weight $w_{k,j}$ has its own rate $\eta_{k,j}$
 - 🐾 If $\Delta w_{k,j}$ remains in the same direction, increase $\eta_{k,j}$ (E has a smooth curve in the vicinity of current w)
 - 🐾 If $\Delta w_{k,j}$ changes the direction, decrease $\eta_{k,j}$ (E has a rough curve in the vicinity of current w)

Overfitting in Neural Networks





Overfitting in Neural Networks



Overfitting in Neural Networks

How to address the overfitting problem

-  Weight decay: decrease each weight by some small factor during each iteration
-  Use a validation set of data

Practical Considerations

- 🐾 A good BP net requires more than the core of the learning algorithms. Many parameters must be carefully selected to ensure a good performance.
- 🐾 Although the deficiencies of BP nets cannot be completely cured, some of them can be eased by some practical means.

🐾 Initial weights (and biases)

- 🐾 Random, [-0.05, 0.05], [-0.1, 0.1], [-1, 1]
- 🐾 Normalize weights for hidden layer ($w^{(1,0)}$) (Nguyen-Widrow)
 - 🐾 Random assign initial weights for all hidden nodes
 - 🐾 For each hidden node j , normalize its weight by

$$w_{j,i}^{(1,0)} = \beta \cdot w_{j,i}^{(1,0)} / \left\| w_j^{(1,0)} \right\|_2 \quad \text{where } \beta = 0.7\sqrt[n]{m}$$

$m = \#$ of hidden nodes, $n = \#$ of input nodes

$$\left\| w_j^{(1,0)} \right\|_2 = \beta \text{ after normalization}$$

- 🐾 Avoid bias in weight initialization:








Practical Considerations

🐾 Training samples:

- 🐾 Quality and quantity of training samples often determines the quality of learning results
- 🐾 Samples must collectively represent well the problem space
 - 🐾 Random sampling
 - 🐾 Proportional sampling (with prior knowledge of the problem space)
- 🐾 # of training patterns needed: There is no theoretically ideal number.
 - 🐾 Baum and Haussler (1989): $P = W/e$, where
 - W: total # of weights to be trained (depends on net structure)
 - e: acceptable classification error rate
 - If the net can be trained to correctly classify $(1 - e/2)P$ of the P training samples, then classification accuracy of this net is $1 - e$ for input patterns drawn from the same sample space
 - Example: $W = 27$, $e = 0.05$, $P = 540$. If we can successfully train the network to correctly classify $(1 - 0.05/2) * 540 = 526$ of the samples, the net will work correctly 95% of time with other input.

Practical Considerations

How many hidden layers and hidden nodes per layer:

-  Theoretically, one hidden layer (possibly with many hidden nodes) is sufficient for any L2 functions
-  There is no theoretical results on minimum necessary # of hidden nodes
-  Practical rule of thumb:
 -  $n = \text{\# of input nodes}$; $m = \text{\# of hidden nodes}$
 -  For binary/bipolar data: $m = 2n$
 -  For real data: $m \gg 2n$
-  Multiple hidden layers with fewer nodes may be trained faster for similar quality in some applications

Practical Considerations

🌸 Data representation:

🌸 Binary vs. bipolar

- 🌸 Bipolar representation uses training samples more efficiently

$$\Delta w_{j,i}^{(1,0)} = \eta \cdot \mu_j \cdot x_i \quad \Delta w_{k,j}^{(2,1)} = \eta \cdot \delta_k \cdot x_j^{(1)}$$

no learning will occur when $x_i = 0$ or $x_j^{(1)} = 0$ with binary rep.

- 🌸 # of patterns can be represented with n input nodes:

binary: 2^n

bipolar: $2^{(n-1)}$ if no biases used, this is due to (anti) symmetry
(if output for input x is o , output for input $-x$ will be $-o$)

🌸 Real value data




- 🌸 Input nodes: real value nodes (may subject to normalization)
- 🌸 Hidden nodes with sigmoid or other non-linear function
- 🌸 Node function for output nodes: often linear (even identity)

e.g.,
$$o_k = \sum w_{k,j}^{(2,1)} x_j^{(1)}$$







- 🌸 Training may be much slower than with binary/bipolar data (some use binary encoding of real values)

Neural Network as a Classifier

Weakness

-  Long training time
-  Require a number of parameters typically best determined empirically, e.g., the network topology or “structure.”
-  Poor interpretability: Difficult to interpret the symbolic meaning behind the learned weights and of “hidden units” in the network

Strength

-  High tolerance to noisy data
-  Ability to classify untrained patterns
-  Well-suited for continuous-valued inputs and outputs
-  Successful on a wide array of real-world data
-  Algorithms are inherently parallel
-  Techniques have recently been developed for the extraction of rules from trained neural networks