

Poster Project

Leonel Salazar, Collin Real, Joaquin Ramirez, Seth Harris

```
library(fmsb)
```

```
Warning: package 'fmsb' was built under R version 4.3.3
```

```
library(gbm)
```

```
Warning: package 'gbm' was built under R version 4.3.3
```

```
Loaded gbm 2.2.2
```

```
This version of gbm is no longer under development. Consider transitioning to  
gbm3, https://github.com/gbm-developers/gbm3
```

```
library(here)
```

```
Warning: package 'here' was built under R version 4.3.3
```

```
here() starts at C:/Users/Leonel/Desktop/MSDA/MS Data Analytics/Current Class
```

```
library(ggplot2)
```

```
Warning: package 'ggplot2' was built under R version 4.3.3
```

```
library(gridExtra)
```

```
library(MASS)
```

```
library(corrplot)
```

```
Warning: package 'corrplot' was built under R version 4.3.3
```

```
corrplot 0.92 loaded
```

```
library(caret)
```

```
Warning: package 'caret' was built under R version 4.3.3
```

```
Loading required package: lattice
```

```
Warning: package 'lattice' was built under R version 4.3.3
```

```
Registered S3 methods overwritten by 'pROC':
```

```
  method      from
```

```
  print.roc fmsb
```

```
  plot.roc  fmsb
```

```
library(e1071)
```

```
Warning: package 'e1071' was built under R version 4.3.3
```

```
library(pROC)
```

Warning: package 'pROC' was built under R version 4.3.3

Type 'citation("pROC")' for a citation.

Attaching package: 'pROC'

The following object is masked from 'package:fmsb':

roc

The following objects are masked from 'package:stats':

cov, smooth, var

```
library(nnet)
```

```
library(pROC)
```

```
library(randomForest)
```

Warning: package 'randomForest' was built under R version 4.3.3

randomForest 4.7-1.1

Type rfNews() to see new features/changes/bug fixes.

Attaching package: 'randomForest'

The following object is masked from 'package:gridExtra':

combine

The following object is masked from 'package:ggplot2':

margin

```
library(class)
```

```
# Use 'here' to create the path to the dataset
```

```
diabetes <- here::here("diabetes.csv")
```

```
# Load the dataset and assign it to the variable 'diabetes'
```

```
diabetes <- read.csv(diabetes)
```

```
# View the first few rows of the dataset
```

```
head(diabetes)
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI
1	6	148	72	35	0	33.6
2	1	85	66	29	0	26.6

```

3           8      183           64           0           0 23.3
4           1       89           66          23          94 28.1
5           0      137           40          35         168 43.1
6           5      116           74           0           0 25.6

```

```

DiabetesPedigreeFunction Age Outcome
1           0.627  50      1
2           0.351  31      0
3           0.672  32      1
4           0.167  21      0
5           2.288  33      1
6           0.201  30      0

```

```

# Convert the outcome variable to a factor
diabetes$Outcome <- as.factor(diabetes$Outcome)

```

```

numeric_columns <- names(diabetes)[names(diabetes) != "Outcome"]
diabetes[numeric_columns] <- lapply(diabetes[numeric_columns], as.numeric)

```

```
str(diabetes)
```

```

'data.frame':  768 obs. of  9 variables:
 $ Pregnancies      : num  6 1 8 1 0 5 3 10 2 8 ...
 $ Glucose          : num  148 85 183 89 137 116 78 115 197 125 ...
 $ BloodPressure    : num  72 66 64 66 40 74 50 0 70 96 ...
 $ SkinThickness    : num  35 29 0 23 35 0 32 0 45 0 ...
 $ Insulin          : num  0 0 0 94 168 0 88 0 543 0 ...
 $ BMI              : num  33.6 26.6 23.3 28.1 43.1 25.6 31 35.3 30.5
 0 ...
 $ DiabetesPedigreeFunction: num  0.627 0.351 0.672 0.167 2.288 ...
 $ Age              : num  50 31 32 21 33 30 26 29 53 54 ...
 $ Outcome          : Factor w/ 2 levels "0","1": 2 1 2 1 2 1 2 1 2 2
 ...

```

```

# Original data histograms
for (col in numeric_columns) {
  print(ggplot(diabetes, aes_string(x = col)) +
    geom_histogram(bins = 30, fill = "blue", color = "black", alpha =
0.7) +
    ggtitle(paste("Histogram of", col, "(Original)")) +
    xlab(col) +
    ylab("Frequency"))
}

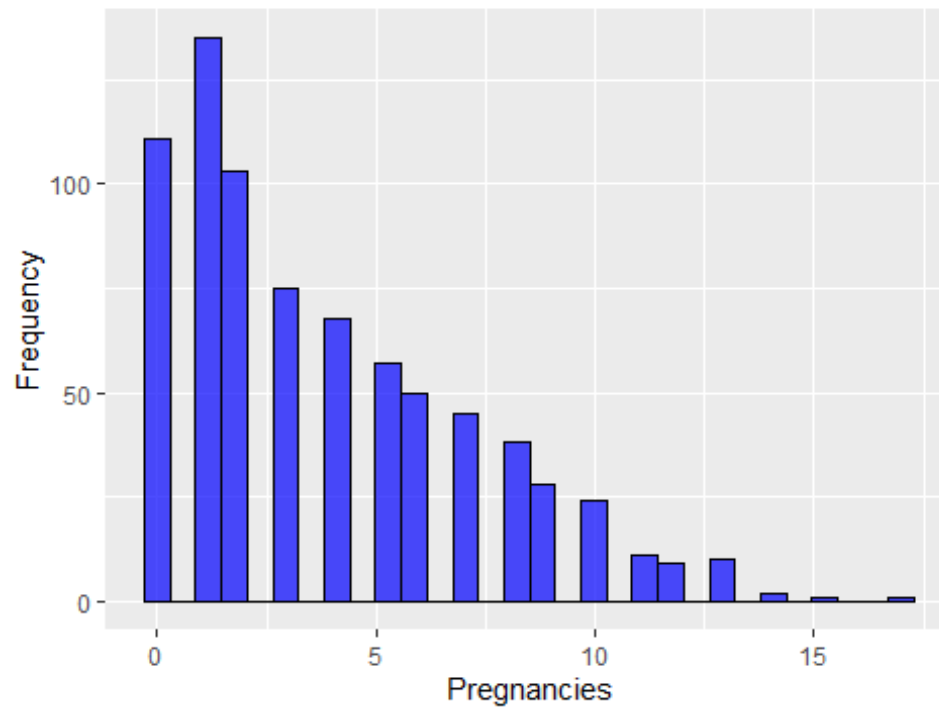
```

```


Warning: `aes_string()` was deprecated in ggplot2 3.0.0.
i Please use tidy evaluation idioms with `aes()`.
i See also `vignette("ggplot2-in-packages")` for more information.

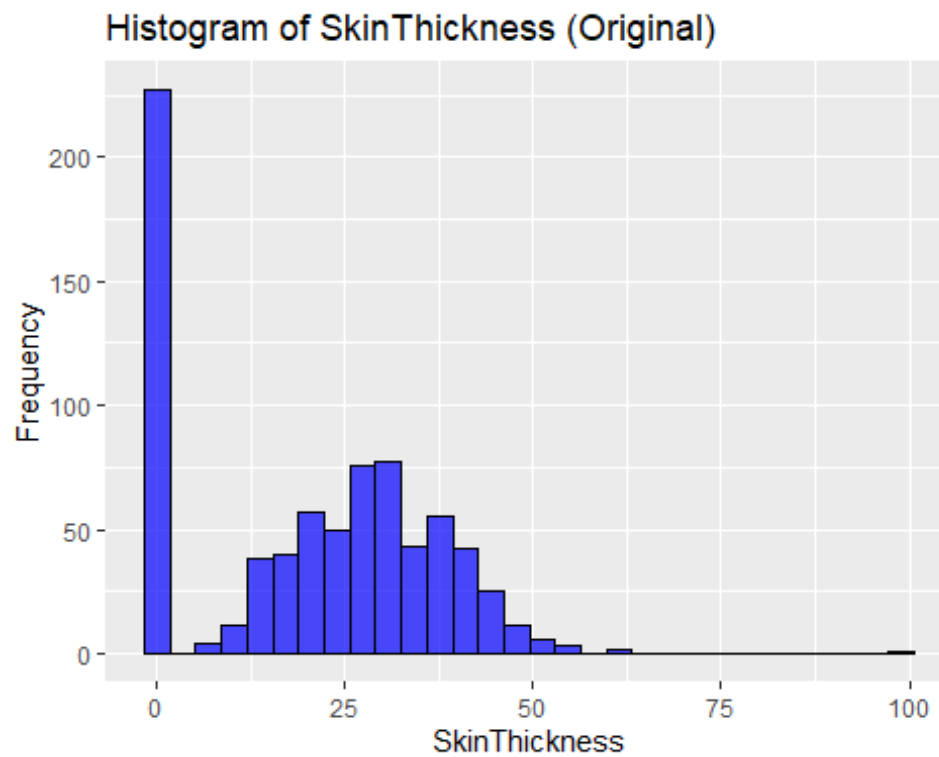

```

Histogram of Pregnancies (Original)

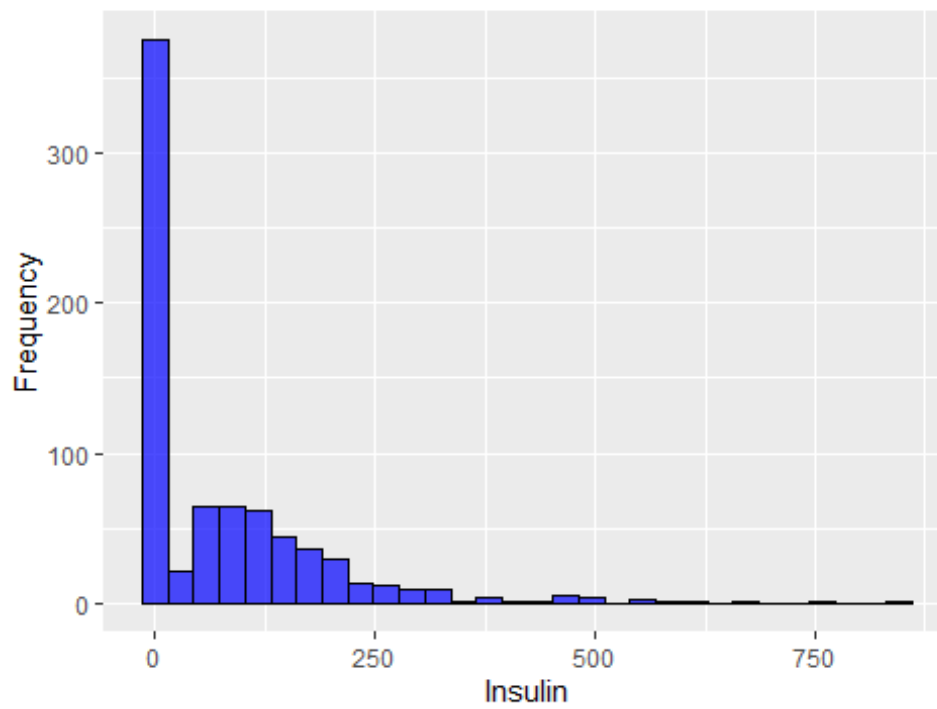


Histogram of Glucose (Original)








Histogram of Insulin (Original)



Histogram of BMI (Original)






```
# Original data Q-Q plots
for (col in numeric_columns) {
  print(ggplot(diabetes, aes_string(sample = col)) +
```


```

stat_qq() +
stat_qq_line() +
ggtitle(paste("Q-Q Plot of", col, "(Original)")) +
xlab("Theoretical Quantiles") +
ylab("Sample Quantiles")
}


```




Q-Q Plot of Glucose (Original)




Q-Q Plot of BloodPressure (Original)




Q-Q Plot of SkinThickness (Original)




Q-Q Plot of Insulin (Original)



Q-Q Plot of BMI (Original)



Q-Q Plot of DiabetesPedigreeFunction (Original)





```
# Load the dataset
diabetes_file_path <- here::here("diabetes.csv")
diabetes <- read.csv(diabetes_file_path)

# Convert the Outcome variable to a factor
diabetes$Outcome <- as.factor(diabetes$Outcome)

# Set seed for reproducibility
set.seed(123)

# Create training (20%) and testing (80%) indices
train_index <- createDataPartition(diabetes$Outcome, p = 0.2, list = FALSE)

# Split the data
train_data <- diabetes[train_index, ]
test_data <- diabetes[-train_index, ]

# Verify the split
cat("Training set size:", nrow(train_data), "\n")

Training set size: 154

cat("Testing set size:", nrow(test_data), "\n")

Testing set size: 614
```

```

# Normalize the predictor variables for neural network
preproc <- preProcess(train_data[, -which(names(train_data) == "Outcome")],
method = c("center", "scale"))
train_data_norm <- predict(preproc, train_data[, -which(names(train_data) ==
"Outcome")])
test_data_norm <- predict(preproc, test_data[, -which(names(test_data) ==
"Outcome")])

# Add the Outcome variable back to the normalized data
train_data_norm$Outcome <- train_data$Outcome
test_data_norm$Outcome <- test_data$Outcome

# Fit logistic regression model
lr_model <- glm(Outcome ~ ., data = train_data, family = binomial)
lr_prob_predictions <- predict(lr_model, newdata = test_data, type =
"response")

# Evaluate logistic regression model
lr_predictions <- ifelse(lr_prob_predictions > 0.5, 1, 0)
lr_confusion <- confusionMatrix(as.factor(lr_predictions), test_data$Outcome)
lr_accuracy <- lr_confusion$overall["Accuracy"]
lr_precision <- lr_confusion$byClass["Pos Pred Value"]
lr_recall <- lr_confusion$byClass["Sensitivity"]
lr_f1 <- 2 * (lr_precision * lr_recall) / (lr_precision + lr_recall)

# Logistic Regression ROC Curve
lr_roc_curve <- roc(test_data$Outcome, lr_prob_predictions)

Setting levels: control = 0, case = 1

Setting direction: controls < cases

# Fit SVM model
svm_model <- svm(Outcome ~ ., data = train_data, kernel = "radial",
probability = TRUE)
svm_predictions <- predict(svm_model, newdata = test_data, probability =
TRUE)
svm_prob_predictions <- attr(svm_predictions, "probabilities")[, 2]

# Evaluate SVM model
svm_confusion <- confusionMatrix(as.factor(ifelse(svm_prob_predictions > 0.5,
1, 0)), test_data$Outcome)
svm_accuracy <- svm_confusion$overall["Accuracy"]
svm_precision <- svm_confusion$byClass["Pos Pred Value"]
svm_recall <- svm_confusion$byClass["Sensitivity"]
svm_f1 <- 2 * (svm_precision * svm_recall) / (svm_precision + svm_recall)

# SVM ROC Curve
svm_roc_curve <- roc(test_data$Outcome, svm_prob_predictions)

```

Setting levels: control = 0, case = 1

Setting direction: controls < cases

```
# Fit k-NN model (k = 5)
knn_model <- knn(train = train_data[, -which(names(train_data) ==
"Outcome")],
                test = test_data[, -which(names(test_data) == "Outcome")],
                cl = train_data$Outcome, k = 5, prob = TRUE)
knn_prob_predictions <- attr(knn_model, "prob")
```

```
# Evaluate k-NN model
knn_confusion <- confusionMatrix(knn_model, test_data$Outcome)
knn_accuracy <- knn_confusion$overall["Accuracy"]
knn_precision <- knn_confusion$byClass["Pos Pred Value"]
knn_recall <- knn_confusion$byClass["Sensitivity"]
knn_f1 <- 2 * (knn_precision * knn_recall) / (knn_precision + knn_recall)
```

k-NN ROC Curve

```
knn_roc_curve <- roc(test_data$Outcome, knn_prob_predictions)
```

Setting levels: control = 0, case = 1

Setting direction: controls < cases

```
# Fit neural network model
nn_model <- nnet(Outcome ~ ., data = train_data_norm, size = 5, maxit = 200,
decay = 0.01, linout = FALSE)
```

weights: 51

```
initial value 97.168442
iter 10 value 52.844198
iter 20 value 35.472041
iter 30 value 26.606381
iter 40 value 22.522330
iter 50 value 21.913196
iter 60 value 21.767813
iter 70 value 21.758312
iter 80 value 21.753819
iter 90 value 21.753647
final value 21.753645
converged
```

```
nn_prob_predictions <- predict(nn_model, newdata = test_data_norm, type =
"raw")
```

Ensure predictions are factors with the same levels as the actual outcomes

```
nn_predictions <- ifelse(nn_prob_predictions > 0.5, 1, 0)
nn_predictions <- as.factor(nn_predictions)
levels(nn_predictions) <- levels(test_data_norm$Outcome)
```

```

# Evaluate neural network model
nn_confusion <- confusionMatrix(nn_predictions, test_data_norm$Outcome)
nn_accuracy <- nn_confusion$overall["Accuracy"]
nn_precision <- nn_confusion$byClass["Pos Pred Value"]
nn_recall <- nn_confusion$byClass["Sensitivity"]
nn_f1 <- 2 * (nn_precision * nn_recall) / (nn_precision + nn_recall)

# Neural Network ROC Curve
nn_roc_curve <- roc(test_data_norm$Outcome, as.numeric(nn_prob_predictions))

Setting levels: control = 0, case = 1

Setting direction: controls < cases

# Fit random forest model
rf_model <- randomForest(Outcome ~ ., data = train_data, ntree = 100, mtry =
3, importance = TRUE)
rf_prob_predictions <- predict(rf_model, newdata = test_data, type =
"prob")[, 2]

# Evaluate random forest model
rf_confusion <- confusionMatrix(predict(rf_model, newdata = test_data),
test_data$Outcome)
rf_accuracy <- rf_confusion$overall["Accuracy"]
rf_precision <- rf_confusion$byClass["Pos Pred Value"]
rf_recall <- rf_confusion$byClass["Sensitivity"]
rf_f1 <- 2 * (rf_precision * rf_recall) / (rf_precision + rf_recall)

# Random Forest ROC Curve
rf_roc_curve <- roc(test_data$Outcome, rf_prob_predictions)

Setting levels: control = 0, case = 1

Setting direction: controls < cases

# Create a data frame to compare models
model_comparison1 <- data.frame(
  Model = c("Logistic Regression", "SVM", "k-NN", "Neural Network", "Random
Forest"),
  Accuracy = c(lr_accuracy, svm_accuracy, knn_accuracy, nn_accuracy,
rf_accuracy),
  Precision = c(lr_precision, svm_precision, knn_precision, nn_precision,
rf_precision),
  Recall = c(lr_recall, svm_recall, knn_recall, nn_recall, rf_recall),
  F1_Score = c(lr_f1, svm_f1, knn_f1, nn_f1, rf_f1)
)


# Print the comparison table
print(model_comparison1)

```

	Model	Accuracy	Precision	Recall	F1_Score
1	Logistic Regression	0.7524430	0.7743363	0.8750	0.8215962
2	SVM	0.7459283	0.7735426	0.8625	0.8156028
3	k-NN	0.6954397	0.7371938	0.8275	0.7797409
4	Neural Network	0.7035831	0.7725000	0.7725	0.7725000
5	Random Forest	0.7345277	0.7749420	0.8350	0.8038508

```
# Plot ROC Curves
```

```
plot(lr_roc_curve, col = "black", main = "ROC Curves for Different Models")
plot(svm_roc_curve, add = TRUE, col = "red")
plot(knn_roc_curve, add = TRUE, col = "blue")
plot(nn_roc_curve, add = TRUE, col = "green")
plot(rf_roc_curve, add = TRUE, col = "purple")
legend("bottomright", legend = c("Logistic Regression", "SVM", "k-NN",
"Neural Network", "Random Forest"),
      col = c("black", "red", "blue", "green", "purple"), lwd = 2)
```



```
# Print AUC values
```

```
cat("AUC for Logistic Regression:", auc(lr_roc_curve), "\n")
```

```
AUC for Logistic Regression: 0.7883411
```

```
cat("AUC for SVM:", auc(svm_roc_curve), "\n")
```

```
AUC for SVM: 0.8054673
```

```
cat("AUC for k-NN:", auc(knn_roc_curve), "\n")
```


AUC for k-NN: 0.3767407

```
cat("AUC for Neural Network:", auc(nn_roc_curve), "\n")
```

AUC for Neural Network: 0.7585864

```
cat("AUC for Random Forest:", auc(rf_roc_curve), "\n")
```

AUC for Random Forest: 0.7944042

Removed all Zeroes

```
library(here)
library(caret)
library(e1071)
library(class)
library(nnet)
library(randomForest)
library(pROC)
library(ggplot2)
library(gridExtra)

# Load the dataset
diabetes_file_path <- here::here("diabetes.csv")
diabetes <- read.csv(diabetes_file_path)

# Convert the Outcome variable to a factor
diabetes$Outcome <- as.factor(diabetes$Outcome)

# Identify the columns to exclude from zero-checking
exclude_columns <- c("Outcome")

# Check how many rows have a 0 in any column except the Outcome column
rows_with_zero <- apply(diabetes[, !names(diabetes) %in% exclude_columns], 1,
function(row) any(row == 0))

# Remove rows with a 0 in any column except the Outcome column
diabetes_cleaned <- diabetes[!rows_with_zero, ]

# Display the number of remaining observations
cat("Number of observations after removing rows with 0 in any column
(excluding Outcome):", nrow(diabetes_cleaned), "\n")

Number of observations after removing rows with 0 in any column (excluding
Outcome): 336

# Set seed for reproducibility
set.seed(123)

# Create training (80%) and testing (20%) indices
train_index <- createDataPartition(diabetes_cleaned$Outcome, p = 0.8, list =
FALSE)
```

```

# Split the data
train_data <- diabetes_cleaned[train_index, ]
test_data <- diabetes_cleaned[-train_index, ]

# Verify the split
cat("Training set size:", nrow(train_data), "\n")

Training set size: 269

cat("Testing set size:", nrow(test_data), "\n")

Testing set size: 67

# Normalize the predictor variables for neural network and k-NN
preproc <- preProcess(train_data[, -which(names(train_data) == "Outcome")],
method = c("center", "scale"))
train_data_norm <- predict(preproc, train_data[, -which(names(train_data) ==
"Outcome")])
test_data_norm <- predict(preproc, test_data[, -which(names(test_data) ==
"Outcome")])
diabetes_norm <- predict(preproc, diabetes_cleaned[, -
which(names(diabetes_cleaned) == "Outcome")])

# Add the Outcome variable back to the normalized data
train_data_norm$Outcome <- train_data$Outcome
test_data_norm$Outcome <- test_data$Outcome
diabetes_norm$Outcome <- diabetes_cleaned$Outcome

# Fit logistic regression model
lr_model <- glm(Outcome ~ ., data = train_data, family = binomial)
lr_prob_predictions <- predict(lr_model, newdata = test_data, type =
"response")
lr_predictions <- as.factor(ifelse(lr_prob_predictions > 0.5, 1, 0))

# Fit SVM model
svm_model <- svm(Outcome ~ ., data = train_data, kernel = "radial",
probability = TRUE)
svm_prob_predictions <- predict(svm_model, newdata = test_data, probability =
TRUE)
svm_prob_predictions <- attr(svm_prob_predictions, "probabilities")[, 2]
svm_predictions <- as.factor(ifelse(svm_prob_predictions > 0.5, 1, 0))

# Fit k-NN model (k = 5)
knn_model <- knn(train = train_data_norm[, -which(names(train_data_norm) ==
"Outcome")],
test = test_data_norm[, -which(names(test_data_norm) ==
"Outcome")],
cl = train_data_norm$Outcome, k = 5)
knn_predictions <- knn_model

```

```

# Fit neural network model
nn_model <- nnet(Outcome ~ ., data = train_data_norm, size = 5, maxit = 200,
decay = 0.01, linout = FALSE)

# weights: 51
initial value 163.827891
iter 10 value 106.844055
iter 20 value 88.132080
iter 30 value 82.934518
iter 40 value 80.964255
iter 50 value 77.907908
iter 60 value 74.862650
iter 70 value 73.097675
iter 80 value 72.205309
iter 90 value 71.664927
iter 100 value 71.391457
iter 110 value 71.354651
iter 120 value 71.342521
iter 130 value 71.338160
iter 140 value 71.337916
final value 71.337895
converged

nn_prob_predictions <- predict(nn_model, newdata = test_data_norm, type =
"raw")
nn_predictions <- as.factor(ifelse(nn_prob_predictions > 0.5, 1, 0))

# Fit random forest model
rf_model <- randomForest(Outcome ~ ., data = train_data, ntree = 100, mtry =
3, importance = TRUE)
rf_prob_predictions <- predict(rf_model, newdata = test_data, type =
"prob")[, 2]
rf_predictions <- as.factor(ifelse(rf_prob_predictions > 0.5, 1, 0))

# Evaluate Models
evaluate_model <- function(predictions, prob_predictions, true_labels) {
  confusion <- confusionMatrix(predictions, true_labels)
  accuracy <- confusion$overall["Accuracy"]
  precision <- confusion$byClass["Pos Pred Value"]
  recall <- confusion$byClass["Sensitivity"]
  f1 <- 2 * (precision * recall) / (precision + recall)
  auc_val <- auc(roc(true_labels, prob_predictions))
  list(confusion = confusion, accuracy = accuracy, precision = precision,
recall = recall, f1 = f1, auc = auc_val)
}

lr_eval <- evaluate_model(lr_predictions, lr_prob_predictions,
test_data$Outcome)

```

```

Setting levels: control = 0, case = 1
Setting direction: controls < cases

svm_eval <- evaluate_model(svm_predictions, svm_prob_predictions,
test_data$Outcome)

Setting levels: control = 0, case = 1
Setting direction: controls < cases

knn_eval <- evaluate_model(knn_predictions, as.numeric(knn_predictions),
test_data$Outcome)

Setting levels: control = 0, case = 1
Setting direction: controls < cases

nn_eval <- evaluate_model(nn_predictions, as.numeric(nn_prob_predictions),
test_data$Outcome)

Setting levels: control = 0, case = 1
Setting direction: controls < cases

rf_eval <- evaluate_model(rf_predictions, rf_prob_predictions,
test_data$Outcome)

Setting levels: control = 0, case = 1
Setting direction: controls < cases

# Create a data frame to compare models
model_comparison <- data.frame(
  Model = c("Logistic Regression", "SVM", "k-NN", "Neural Network", "Random
Forest"),
  Accuracy = c(lr_eval$accuracy, svm_eval$accuracy, knn_eval$accuracy,
nn_eval$accuracy, rf_eval$accuracy),
  Precision = c(lr_eval$precision, svm_eval$precision, knn_eval$precision,
nn_eval$precision, rf_eval$precision),
  Recall = c(lr_eval$recall, svm_eval$recall, knn_eval$recall,
nn_eval$recall, rf_eval$recall),
  F1_Score = c(lr_eval$f1, svm_eval$f1, knn_eval$f1, nn_eval$f1, rf_eval$f1),
  AUC = c(lr_eval$auc, svm_eval$auc, knn_eval$auc, nn_eval$auc, rf_eval$auc)
)

# Print the comparison table
print(model_comparison)

```

	Model	Accuracy	Precision	Recall	F1_Score	AUC
1	Logistic Regression	0.7910448	0.8039216	0.9111111	0.8541667	0.8434343
2	SVM	0.7611940	0.7843137	0.8888889	0.8333333	0.8292929
3	k-NN	0.7910448	0.8163265	0.8888889	0.8510638	0.7398990
4	Neural Network	0.6567164	0.7391304	0.7555556	0.7472527	0.7767677
5	Random Forest	0.7910448	0.8297872	0.8666667	0.8478261	0.8671717

```

# Extract ROC data
extract_roc_data <- function(roc_object) {
  data.frame(
    specificity = roc_object$specificities,
    sensitivity = roc_object$sensitivities,
    model = deparse(substitute(roc_object))
  )
}

lr_roc_data <- extract_roc_data(roc(test_data$Outcome, lr_prob_predictions))

Setting levels: control = 0, case = 1
Setting direction: controls < cases

svm_roc_data <- extract_roc_data(roc(test_data$Outcome,
svm_prob_predictions))

Setting levels: control = 0, case = 1
Setting direction: controls < cases

knn_roc_data <- extract_roc_data(roc(test_data$Outcome,
as.numeric(knn_predictions)))

Setting levels: control = 0, case = 1
Setting direction: controls < cases

nn_roc_data <- extract_roc_data(roc(test_data$Outcome,
as.numeric(nn_prob_predictions)))

Setting levels: control = 0, case = 1
Setting direction: controls < cases

rf_roc_data <- extract_roc_data(roc(test_data$Outcome, rf_prob_predictions))

Setting levels: control = 0, case = 1
Setting direction: controls < cases

# Combine ROC data
roc_data <- rbind(
  transform(lr_roc_data, model = "Logistic Regression"),
  transform(svm_roc_data, model = "SVM"),
  transform(knn_roc_data, model = "k-NN"),
  transform(nn_roc_data, model = "Neural Network"),
  transform(rf_roc_data, model = "Random Forest")
)


# Plot ROC Curves
roc_plot <- ggplot(roc_data, aes(x = 1 - specificity, y = sensitivity, color
= model)) +
  geom_line(size = 1) +
  labs(x = "False Positive Rate", y = "True Positive Rate", title = "ROC

```

```
Curves") +  
  theme_minimal()
```

Warning: Using `size` aesthetic for lines was deprecated in ggplot2 3.4.0.
 i Please use `linewidth` instead.

```
# Display the ROC plot  
print(roc_plot)
```



```
# Generate predictions for the entire dataset  
lr_prob_predictions <- predict(lr_model, newdata = diabetes_cleaned, type =  
  "response")  
lr_predictions <- as.factor(ifelse(lr_prob_predictions > 0.5, 1, 0))  
  
svm_prob_predictions <- predict(svm_model, newdata = diabetes_cleaned,  
  probability = TRUE)  
svm_prob_predictions <- attr(svm_prob_predictions, "probabilities")[, 2]  
svm_predictions <- as.factor(ifelse(svm_prob_predictions > 0.5, 1, 0))  
  
knn_predictions <- knn(train = train_data_norm[, -  
  which(names(train_data_norm) == "Outcome")],  
  test = diabetes_norm[, -which(names(diabetes_norm) ==  
  "Outcome")],  
  cl = train_data_norm$Outcome, k = 5)  
  
nn_prob_predictions <- predict(nn_model, newdata = diabetes_norm, type =  
  "raw")
```

```

nn_predictions <- as.factor(ifelse(nn_prob_predictions > 0.5, 1, 0))

rf_prob_predictions <- predict(rf_model, newdata = diabetes_cleaned, type =
"prob")[, 2]
rf_predictions <- as.factor(ifelse(rf_prob_predictions > 0.5, 1, 0))

# Create a new dataset with original outcomes and predictions
diabetes_pred <- diabetes_cleaned
diabetes_pred$LR_Prediction <- lr_predictions
diabetes_pred$SVM_Prediction <- svm_predictions
diabetes_pred$KNN_Prediction <- knn_predictions
diabetes_pred$NN_Prediction <- nn_predictions
diabetes_pred$RF_Prediction <- rf_predictions

# Display the first few rows of the new dataset
head(diabetes_pred)

```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI
4	1	89	66	23	94	28.1
7	3	78	50	32	88	31.0
9	2	197	70	45	543	30.5
14	1	189	60	23	846	30.1
15	5	166	72	19	175	25.8
19	1	103	30	38	83	43.3

	DiabetesPedigreeFunction	Age	Outcome	LR_Prediction	SVM_Prediction
4	0.167	21	0	0	0
7	0.248	26	1	0	0
9	0.158	53	1	1	1
14	0.398	59	1	1	1
15	0.587	51	1	1	1
19	0.183	33	0	0	0

	KNN_Prediction	NN_Prediction	RF_Prediction
4	0	0	0
7	0	0	1
9	1	1	1
14	1	1	1
15	1	1	1
19	0	0	0

```

library(ggplot2)


# Create a copy of the model_comparison dataframe
model_comparison_ordered <- model_comparison

# Add a column for the overall ranking based on AUC
model_comparison_ordered$Rank <- rank(-model_comparison_ordered$AUC)

# Sort the dataframe by the ranking
model_comparison_ordered <-
model_comparison_ordered[order(model_comparison_ordered$Rank), ]

```

```
# Plot the rankings
ggplot(model_comparison_ordered, aes(x = reorder(Model, -AUC), y = AUC, fill
= Model)) +
  geom_bar(stat = "identity") +
  labs(title = "Model Comparison: AUC Ranking", x = "Model", y = "AUC") +
  theme_minimal() +
  theme(legend.position = "none")
```



```
library(gridExtra)

# Create individual plots
plot1 <- ggplot(model_comparison, aes(x = Model, y = Accuracy, fill = Model))
+
  geom_bar(stat = "identity") +
  labs(title = "Model Comparison: Accuracy", x = "Model", y = "Accuracy") +
  theme_minimal() +
  theme(legend.position = "none")

plot2 <- ggplot(model_comparison, aes(x = Model, y = Precision, fill =
Model)) +
  geom_bar(stat = "identity") +
  labs(title = "Model Comparison: Precision", x = "Model", y = "Precision") +
  theme_minimal() +
  theme(legend.position = "none")
```



```

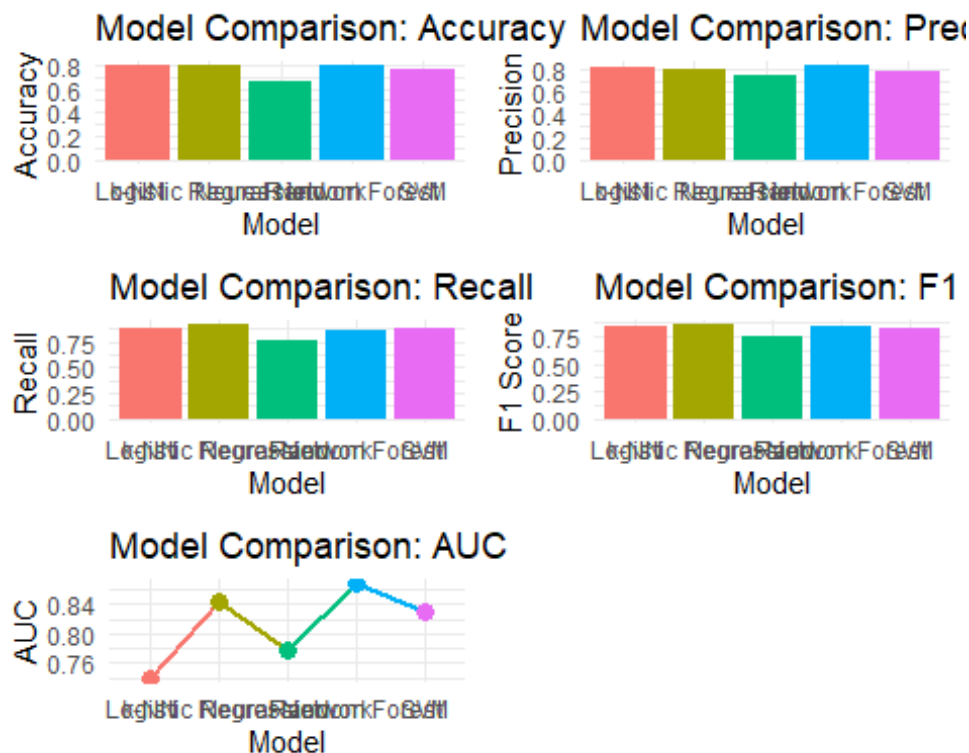
plot3 <- ggplot(model_comparison, aes(x = Model, y = Recall, fill = Model)) +
  geom_bar(stat = "identity") +
  labs(title = "Model Comparison: Recall", x = "Model", y = "Recall") +
  theme_minimal() +
  theme(legend.position = "none")

plot4 <- ggplot(model_comparison, aes(x = Model, y = F1_Score, fill = Model)) +
  geom_bar(stat = "identity") +
  labs(title = "Model Comparison: F1 Score", x = "Model", y = "F1 Score") +
  theme_minimal() +
  theme(legend.position = "none")


plot5 <- ggplot(model_comparison, aes(x = Model, y = AUC, group = 1)) +
  geom_line(aes(color = Model), size = 1) +
  geom_point(aes(color = Model), size = 3) +
  labs(title = "Model Comparison: AUC", x = "Model", y = "AUC") +
  theme_minimal() +
  theme(legend.position = "none")

# Arrange the plots in a grid
grid.arrange(plot1, plot2, plot3, plot4, plot5, ncol = 2)


```




```
print(plot1)
```




```
print(plot2)
```




```
print(plot3)
```



```
print(plot4)
```



```
print(plot5)
```



```
# Fit random forest model
rf_model <- randomForest(Outcome ~ ., data = train_data, ntree = 100, mtry =
3, importance = TRUE)

# Get feature importance
importance <- importance(rf_model)
importance_df <- data.frame(Feature = row.names(importance), Importance =
importance[, "MeanDecreaseGini"])

# Sort by importance
importance_df <- importance_df[order(importance_df$Importance, decreasing =
TRUE), ]

# Print feature importance
print(importance_df)
```

	Feature	Importance
Glucose	Glucose	29.604150
Age	Age	18.760203
Insulin	Insulin	18.053296
DiabetesPedigreeFunction	DiabetesPedigreeFunction	12.559633
SkinThickness	SkinThickness	10.850048
BloodPressure	BloodPressure	10.329266
BMI	BMI	10.173463
Pregnancies	Pregnancies	8.348566

```
ggplot(importance_df, aes(x = reorder(Feature, Importance), y = Importance,
fill = Importance)) +
  geom_bar(stat = "identity") +
  coord_flip() +
  labs(title = "Feature Importance from Random Forest Model", x = "Features",
y = "Importance") +
  theme_minimal() +
  scale_fill_gradient(low = "violet", high = "orange")
```

